

Mattias Villani

Bayesian Learning

A GENTLE INTRODUCTION

Some publisher

Copyright © 2020 Mattias Villani

PUBLISHED BY SOME PUBLISHER

TYPESET BY \LaTeX USING TEMPLATE FROM [TUFTE-LATEX.GITHUB.IO](https://github.com/tufte/tufte-latex)

Licensed under the Apache License, Version 2.0 (the “License”); you may not use this file except in compliance with the License.

First edition, November 2020

Contents

<i>The Bayesics</i>	7
<i>Single-parameter models</i>	17
<i>Multi-parameter models</i>	37
<i>Regression</i>	51
<i>Prediction and Decision making</i>	55
<i>Classification</i>	57
<i>Posterior simulation</i>	59
<i>Variational inference</i>	61
<i>Regularization</i>	63
<i>Model comparison</i>	65
<i>Variable selection</i>	67
<i>Gaussian processes</i>	69
<i>Mixture models</i>	71
<i>Bibliography</i>	73
<i>Index</i>	75

Preface

Who is this book for?

This book can be used as a first book in Bayesian statistics at the advanced undergraduate or master level. The book is written so that it can accomodate also students in engineering and computer science who are interested in Bayesian learning for applications in the field of Machine Learning.

In fact, the book grew out of a Bayesian course that I taught for groups of heterogenous students with roughly half of students from statistics and the other half from engineering and computer science, often with an interest in machine learning. To my surprise, I found that it was indeed possible to teach the same material to all students, even if half the class had a much more extensive background in statistics. The course had always very favorable reviews from the students and not a single student has complained over the years on it being too easy or too hard. There are two main explanations for this. First, since most bachelor level Statistics are non-Bayesian in methods and thinking, taking a first course in Bayesian inference is in some way like starting from scratch. Sure, there are several overlapping concepts and probability is of course the underlying technical language (although with highly different interpretations), but there are nevertheless a lot of effort spend in basic statistics courses that are not needed prerequisites for a Bayesian course. Second, my courses are very computational, as is most of the Bayesian field, with a lot of computer labs and also a partly computerized exam. Engineering and particularly computer science students tend to have a comparative advantage in computing and programming. So the additional time that students from statistics had to spend on programming, computer science students could spend on catching up on statistical concepts. In the end, everyone seemed put in the same number of hours and everyone was happy with the learning experience. In order to accomodate both groups of students, my lectures covers also some rather elementary concepts, especially in the early part of the course, but then rather quickly moves over to territory unknown to all stu-

dents. This book is written in the same style using Tufte style margin notes and figures to fill in potential missing gaps in probability and statistics, without breaking the flow of the main text.

Programming is useful for the exercises, or at least basic familiarity with R, Python or Julia or a similar datacentric language. I will use pseudo code for certain smaller algorithms and Julia for real code; Julia is used to present algorithms in the book since the ability to use mathematical symbols in Julia (via unicode) makes the code easy to read, almost like pseudo code.

Why the term Bayesian learning?

I have used the term Bayesian *learning* in the book's title instead of Bayesian *inference* or Bayesian *statistics*. There are several reasons for this.

First, I want my courses and this book to be welcoming to students in fields neighboring statistics, such as machine learning, computer science, and parts of engineering. This reflects my strong belief that a modern statistician or machine learner should be a little of a renaissance person that understands both probability and statistical modelling, and computing. The ideal class is therefore a mix of students from nearby disciplines that learn for each others competences as much as they learn from my classes or this book.

Second, the term learning instead of inference was chosen since Bayesian statistics is about learning from data, often in a very sequential way where incrementally collected information updates our knowledge about the world.

Finally, the title is meant to convey the message that this is not a traditional book in statistics. The approach taken here, especially in later chapters, is very computationally driven with many algorithms for real-world data analysis. It is also inspired by machine learning in that much of the focus is given to prediction and decision making, and almost none to hypothesis testing.

The Bayesics

Learning probability models

A central task in statistics is to infer an unknown parameter $\theta \in \Theta$ in a probability model $p(X_1, \dots, X_n | \theta)$ from a dataset of n observations x_1, \dots, x_n . The **parameter space** Θ is the set of allowed parameter values. Some examples of problems with a single parameter are learning the voting share of a political party from exit polls, predicting the number of bugs in a software release or inferring a one-dimensional measure of a student's true mathematical ability.

Most problems require models with more than one parameter however. A prominent example with an extremely large number of parameters are the deep neural network models widely used in artificial intelligence (AI); such models often have millions of network weights that have to be learned from training data. However, to focus on ideas and easy derivations, we will keep things as simple as possible in the first two chapters and only consider models with a single parameter. Later chapters tackle more complex models and presents methods specifically designed for models with many parameters.

The initial chapters will be very focused on learning parameters in models. It is important to remember that parameter inference is usually an intermediate step toward the final aim of prediction or decision making under uncertainty; for example, the predictions and decisions of a robot are based on a probability model with network weights learned from training data. Throughout this book we will exclusively work with probability models. Probability models have the distinct advantage of giving a precise quantification of uncertainty that can be directly used for decision making in the real world. The Bayesian approach to predictions and decisions will be presented in the chapter [Prediction and Decision making](#).

Statistical distributions and independent data

We will initially assume that the observations X_1, \dots, X_n are *independent and identically distributed* (**iid**) conditional on θ so that we can

parameter space



Figure 1: Artificial intelligence is often based on Bayesian learning.

iid

write the joint distribution

$$p(X_1, \dots, X_n | \theta) = \prod_{i=1}^n p(X_i | \theta).$$

We will write $X_1, \dots, X_n | \theta \stackrel{\text{iid}}{\sim} p(X | \theta)$ to denote this. In this setting we can refer to ‘the probability model’ as the probability distribution $p(X | \theta)$ for a single observation.

EXAMPLE: A binary random variable $X \in \{0, 1\}$ follows a **Bernoulli distribution** if

$$\Pr(X = x | \theta) = \begin{cases} \theta & \text{for } x = 1 \\ 1 - \theta & \text{for } x = 0 \end{cases}$$

which can be written more compactly as

$$\Pr(X = x | \theta) = \theta^x (1 - \theta)^{1-x}. \quad (1)$$

A typical example of iid Bernoulli data occurs when a coin is flipped n times (also called **Bernoulli trials**) and the sequence of heads ($x = 1$) and tails ($x = 0$) are recorded. It is common to refer to the outcome $X = 1$ as a success, and $X = 0$ as a failure. The Bernoulli distribution is illustrated in Figure 2.

We make the usual distinction between *random variables* denoted by capital letters and their *realizations (data)*, so $X = x$ means a random variable X with outcome x . As we will see later on, this distinction will often be less relevant in a Bayesian world where all inferences are conditioned on the observed data; we will therefore be more sloppy with this distinction in later chapters, but no harm will come from this.

The likelihood function and maximum likelihood estimation

The likelihood function is a key component of Bayesian learning, and indeed in all of Statistics. Given a probability model $p(X_1, \dots, X_n | \theta)$ the **likelihood function** $p(x_1, \dots, x_n | \theta)$ is the *joint* probability of observing the data set x_1, \dots, x_n *considered as a function* of the parameter θ . Again, if the data are iid we can express the likelihood in terms of the univariate distributions $p(X | \theta)$ as

$$p(x_1, \dots, x_n | \theta) = \prod_{i=1}^n p(x_i | \theta). \quad (2)$$

EXAMPLE: In the case of iid Bernoulli data the likelihood function is simply obtained by multiplying together the probability of success

Bernoulli distribution




Figure 2: Bernoulli distribution with success probability $\theta = 0.8$.

Bernoulli trials

likelihood function

θ for the observations where $x_i = 1$ and probability of failure $1 - \theta$ when $x_i = 0$, giving the likelihood

$$p(x_1, \dots, x_n | \theta) = \theta^s (1 - \theta)^f, \quad (3)$$

where $s = \sum_{i=1}^n x_i$ is the number of successes in the sample, and $f = n - s$ is the number of failures.

It is absolutely essential to have mental image of the likelihood function when thinking about statistical modeling. Figure 3 illustrates the likelihood function for Bernoulli model when $s = 4$ successes was obtained in $n = 10$ trials. Figure 4 is another example with the same proportion of 40% successes, but here in $n = 100$ trials; note how the large dataset makes the likelihood more concentrated, more informative regarding the plausibility of different θ values. Finally, Figure 5 is again the likelihood from only $n = 10$ trials, but now with $s = 9$ successes; this data is much more likely to have resulted from a data generating process with θ closer to one.

Figures 3-5 nicely illustrates how the likelihood function can inform us about the plausibility of different θ for any given dataset. If we want to select a single value, an estimate of θ , a natural candidate is the widely used **maximum likelihood estimator**

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta \in \Theta} p(x_1, \dots, x_n | \theta). \quad (4)$$

It makes some intuitive sense to estimate θ by the value that maximizes the probability of the observed data; the estimator $\hat{\theta}_{\text{MLE}}$ also enjoys several other attractive properties, particularly in large samples, i.e. when n is large.

It is quite easy to derive $\hat{\theta}_{\text{MLE}}$ for iid Bernoulli data. Rather than maximizing $p(x_1, \dots, x_n | \theta)$ directly with respect to θ it is often easier to maximize the *log-likelihood function*

$$\log p(x_1, \dots, x_n | \theta) = s \log \theta + f \log(1 - \theta).$$

Since the logarithm is a monotonically increasing function we obtain the same estimator if we maximize the likelihood or the log-likelihood function. We can now easily find $\hat{\theta}_{\text{MLE}}$ by taking the first derivative of the log-likelihood function with respect to θ , setting that derivative to zero and solving for θ . Solving

$$\frac{d \log p(x_1, \dots, x_n | \theta)}{d\theta} = \frac{s}{\theta} - \frac{f}{1 - \theta} = 0,$$

gives the unique solution $\hat{\theta}_{\text{MLE}} = s/n$, the fraction of successes in the data. It is straightforward to show that this indeed a maximum by checking that the second derivative is negative at $\hat{\theta}_{\text{MLE}}$.

The maximum likelihood estimator is **unbiased** in this example,




Figure 3: Bernoulli likelihood function for $n = 10$ and $s = 4$.




Figure 4: Bernoulli likelihood function for $n = 100$ and $s = 40$.

maximum likelihood estimator




Figure 5: Bernoulli likelihood function for $n = 10$ and $s = 9$.

Binomial distribution

$S \sim \text{Binom}(n, \theta)$
Support: $S \in 0, 1, \dots, n$

$$p(s) = \binom{n}{s} \theta^s (1 - \theta)^{n-s}$$

$$\mathbb{E}(X) = n\theta$$

$$\mathbb{V}(X) = n\theta(1 - \theta)$$

Figure 6: The binomial distribution.

unbiased

i.e. it is correct on average over all possible samples from the model:

$$\mathbb{E} [\hat{\theta}_{MLE}(X_1, \dots, X_n)] = \mathbb{E} \left(\frac{S}{n} \right) = \frac{n\theta}{n} = \theta,$$

where we have written out explicitly that an estimator is function of the sample. Note that the number of successes is random in this calculation as we are considering the variability over all possible samples, hence the use of capital letter S . We have also used that if $X_1, \dots, X_n | \theta \stackrel{iid}{\sim} \text{Bernoulli}$ then $S | \theta \sim \text{Binomial}(n, \theta)$ with mean $E(S) = n\theta$; see Figure 7 for an example of a **Binomial distribution**.

The **sampling variance** of an estimator is often used to assess the quality of an estimator. It is easily calculated for $\hat{\theta}_{MLE}$ in the Bernoulli example as

$$\mathbb{V} [\hat{\theta}_{MLE}(X_1, \dots, X_n)] = \mathbb{V} \left(\frac{S}{n} \right) = \frac{1}{n^2} \mathbb{V}(S) = \frac{\theta(1-\theta)}{n},$$

since $\mathbb{V}(S) = n\theta(1-\theta)$ when $S | \theta \sim \text{Binomial}(n, \theta)$.

It is important to understand that the above mean and variance of $\hat{\theta}_{MLE}$ are computed with respect to the **sampling distribution**, i.e. the distribution of the estimator as we repeatedly sample new datasets of size n from the assumed data generating process. They are long run properties of the estimation method, telling us how the estimator would perform on average over many repeatedly sampled datasets. Such long run properties play a very limited role in the Bayesian approach to inference where one can directly condition the inferences on the single dataset that we have observed. While the sampling properties of $\hat{\theta}_{MLE}$ are not used in Bayesian world, the likelihood *function* is at the core of Bayesian learning.

The likelihood functions in Figures 3-5 look like a probability distribution for θ , and it is tempting to compute probabilities for θ , for example $\Pr(\theta \leq c | x_1, \dots, x_n)$ for some c . Of course, such probabilities only makes sense if θ is a random variable, and we have so far considered θ to be a fixed unknown constant. So while $p(X_1, \dots, X_n | \theta)$ is a probability distribution for a random sample X_1, \dots, X_n for a fixed θ , the likelihood function is only the probability of a *fixed* sample x_1, \dots, x_n considered as function of θ ; the likelihood is therefore *not* a probability distribution for θ . Figure 8 reminds us of this error.

This is somewhat disappointing since having a probability distribution for θ would be very useful, for example when making a decision whose consequences depend on the unknown θ ; see Chapter **Prediction and Decision making**. But again, it only makes sense to speak about probabilities for θ when θ is random. And this is where our Bayesian story begins.

Binomial distribution




Figure 7: Binomial distribution with $n = 10$ and $\theta = 0.7$.

sampling variance

sampling distribution




Figure 8: The likelihood function is **not** a probability distribution for θ .

Subjective Probability

What is the probability that the 10th decimal of π is 3? This may seem like a silly question since there is nothing intrinsically random about the 10th decimal of π ; it is a fixed quantity that does not vary. A Bayesian will however argue that if *you do not know its value* then you should express that uncertainty by a probability distribution. The Italian mathematician Bruno de Finetti, one of the founders of this school of probability, has expressed this well:

The only relevant thing is uncertainty - the extent of our knowledge and ignorance. The actual fact of whether or not the events considered are in some sense determined, or known by other people, and so on, is of no consequence.

Bruno de Finetti in his 1974 book 'A Theory of Probability' Vol 1.

Probability is the language of uncertainty and Bayesian learning is based on a **subjective probability**. A subjective probability measures the **personal degree of belief** of a person. Since different person have difference knowledge and experience, such beliefs will vary between persons. A person that has no idea about the 10th decimal of π may use a uniform distribution on the integers 0-9. Someone else may however know this decimal with certainty and assigns a probability of 1 to that outcome. Again, whether or not the event is in some sense intrinsically random or not is of no consequence; the only relevant thing is *your* uncertainty. Einstein's famous statement "God does not play dice with the universe" may interesting to ponder about, but has not bearing on subjective probability and Bayesian learning.

The notion of probability in Bayesian learning is therefore radically different from the frequentist interpretation of probability taught in most basic statistics classes. The **frequentist probability** of an event A is defined as the limiting proportion of times that event A occurs in an (imagined) infinite number of repetitions of an experiment; for example the tossing a coin with the event of interest $A = \{\text{Heads}\}$. A subjective probability measure is instead defined as the personal degree of belief in the event A for a person. Note that subjective probabilities can be used to quantify uncertainties also for events that are unrepeatable, for example the probability of a nuclear disaster at a particular location. A subjective probability distribution can also contain useful information that may not directly come from observed data. As we will see, the Bayesian approach combines such subjective information with objective data in a natural way.

Luckily, the computational rules for probabilities are the same for both frequentist and subjective interpretations of probability; for

π

subjective probability

personal degree of belief



Figure 9: Bruno de Finetti, 1906-1985, a founder of subjective probability.

frequentist probability

example $0 \leq \Pr(A) \leq 1$ and $\Pr(A \cap B) = \Pr(A) + \Pr(B)$ when A and B are disjoint events. The rules can be motivated by considering subjective probabilities as the result of pricing of bets. Imagine that you are given the chance to enter a bet where you win \$1 if event A occurs. How much would you be willing to pay that bet? Surely not more than \$1 as then you would lose money with certainty. If you strongly believe that A will occur you would probably be willing to pay closer to \$1, but if you believe that A is nearly impossible your price for the bet would be close to \$0. The highest price that you would be willing to pay for the bet is your subjective probability in the event A . Given this setup one can easily show that your subjective probabilities must satisfy the axioms for probabilities otherwise you would be willing to enter a sequence of bets where you would lose an infinite amount with certainty; this is the so called **dutch book argument**. Objections have been raised against this argument, for example that the utility from the bet may not linearly increase with the monetary gain, and some people may even get utility just by the excitement in gambling; subsequent refinements of this argument have therefore completely disposed with the notion of money in favor of a more general notion of utility; see the chapter [Prediction and Decision making](#).

dutch book argument

Bayesian Learning

The general recipe for Bayesian learning about an event A is:

- Formulate your subjective *prior beliefs* $\Pr(A)$ about A .
- *Collect data* that inform you about A .
- *Update* your prior beliefs with the observed data.

The big question is *how* to update prior beliefs with data. Bayesian learning gets its name from using Bayes' theorem for this updating. The most basic version of **Bayes' theorem** computes the probability of an event A given the known occurrence of some other event B as

$$\Pr(A|B) = \frac{\Pr(B|A)\Pr(A)}{\Pr(B)}$$

One way to think about this result is that it 'reverses the conditioning', i.e. it computes $\Pr(A|B)$ from $\Pr(B|A)$.

Bayes' theorem will be used to infer an unknown parameter θ in a probability model, but let us first use the theorem to solve a simple problem. Imagine that you have taken a test for a specific latent disease and that the test was unfortunately positive. The doctor tells you that $\Pr(B|A) = 0.9$ where $A = \{\text{'Have disease'}\}$ and

Bayes' theorem



Figure 10: Reverend Thomas Bayes, ca 1701-1761, whose theorem was published posthumously.

$B = \{\text{'Positive test'}\}$ and also $\Pr(B|A^c) = 0.05$, where A^c is the complement to A , i.e. the event that you do not have the disease. Hence, a positive test is very unlikely if you do not have the disease, so you start to worry. But what you really want to know is the probability of having the disease given a positive test, i.e. $\Pr(A|B)$. To compute this you need to know the so called *prior* probability of A before you took the test. The doctor tells you that only one in ten thousand has the disease and you set $\Pr(A) = 0.0001$. Bayes' theorem then gives

$$\Pr(A|B) = \frac{\Pr(B|A)\Pr(A)}{\Pr(B)} = \frac{\Pr(B|A)\Pr(A)}{\Pr(B|A)\Pr(A) + \Pr(B|A^c)\Pr(A^c)} \approx 0.0018,$$

where we have expressed $\Pr(B)$ in the numerator using a version of the **law of total probability**. Hence, even though the test has increased the probability of having the disease by a factor of 18 from the initial $\Pr(A) = 0.0001$, the probability of actually having the disease is still tiny. The lesson here is that prior probabilities matter.

law of total probability

To see how Bayes' theorem can be used for Bayesian learning from data, let us consider the event $B = \{\text{'Data } x_1, \dots, x_n \text{ was observed'}\}$ which we write simply as $B = \{x_1, \dots, x_n\}$. We can now use Bayes' theorem to update the initial beliefs $\Pr(A)$ about some event A with data $B = \{x_1, \dots, x_n\}$ by the formula

$$\Pr(A|x_1, \dots, x_n) = \frac{\Pr(x_1, \dots, x_n|A)\Pr(A)}{\Pr(x_1, \dots, x_n)}.$$

The initial belief $\Pr(A)$ is called a **prior** since it refers to beliefs about A *before* the data x_1, \dots, x_n was observed. Likewise $\Pr(A|x_1, \dots, x_n)$ is referred to as the **posterior** since it is the probability of A *after* data was observed.

prior

posterior

Let us now show how Bayes' theorem can be used to infer a parameter in a probability model $p(X_1, \dots, X_n|\theta)$. We first take a simplified approach where the only possible parameter values are on a grid of values $\theta_1, \theta_2, \dots, \theta_K$. Let $B = \{x_1, \dots, x_n\}$ be the event of observing a specific dataset and $A_k = \{\theta_k\}$ be the event that $\theta = \theta_k$. The posterior probability for each $A_k = \{\theta_k\}$ is then

$$\Pr(\theta_k|x_1, \dots, x_n) = \frac{\Pr(x_1, \dots, x_n|\theta_k)\Pr(\theta_k)}{\sum_{j=1}^K \Pr(x_1, \dots, x_n|\theta_j)\Pr(\theta_j)}. \quad (5)$$

Note how we again used the law of total probability in the denominator to express $\Pr(B) = \Pr(x_1, \dots, x_n)$. This denominator is only there to guarantee that the posterior is a probability distribution, i.e. that $\sum_{j=1}^K \Pr(\theta_j|x_1, \dots, x_n) = 1$.

The really interesting stuff is however in the numerator of (5) and we will therefore often write Bayes' theorem in proportional form

$$\Pr(\theta_k|x_1, \dots, x_n) \propto \Pr(x_1, \dots, x_n|\theta_k)\Pr(\theta_k), \quad (6)$$

where the symbol \propto is read as ‘is proportional to’, i.e. a multiplicative normalizing constant is missing in the expression. Now here is the really crucial thing: the factor $\Pr(x_1, \dots, x_n | \theta_k)$ in Equation (6) is the *likelihood function* evaluated in the point θ_k . Equation (6) therefore expresses the fundamental idea in Bayesian learning:

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior}. \quad (7)$$

Figure 11-13 illustrates the updating from prior to posterior for the Bernoulli model with data $n = 10$ and $s = 9$ over a grid of θ values. Note how the posterior is a compromise between the prior information and the data information (likelihood).

Finally, taking a finer and finer grid in Equation 5 we get the following Bayes’ theorem for a continuous parameter θ in the limit

$$p(\theta | x_1, \dots, x_n) = \frac{p(x_1, \dots, x_n | \theta) p(\theta)}{\int p(x_1, \dots, x_n | \theta) p(\theta) d\theta} \quad (8)$$

where $p(\theta)$ is now a continuous **prior density** that gets updated with new data via the likelihood function $p(x_1, \dots, x_n | \theta)$ to a **posterior density** $p(\theta | x_1, \dots, x_n)$. The normalizing constant is now given by an integral over θ and is a continuous version of the law of total probability. We can again hide the unimportant normalizing constant to get the nicer form

$$p(\theta | x_1, \dots, x_n) \propto p(x_1, \dots, x_n | \theta) p(\theta). \quad (9)$$

It is important to note that the posterior distribution $p(\theta | x_1, \dots, x_n)$ is a probability distribution for the parameter θ ; it completely describes the knowledge about θ for a person with the prior $p(\theta)$ after having observed the data x_1, \dots, x_n . Remember that the likelihood can not be used to compute probabilities for θ . With a posterior distribution we actually *can* compute $\Pr(\theta \leq c | x_1, \dots, x_n) = \int p(\theta \leq c | x_1, \dots, x_n) d\theta$ or any other probability of interest. It is the prior $p(\theta)$ that makes it possible to use Bayes’ theorem to revert the conditioning in the likelihood $p(x_1, \dots, x_n | \theta)$ into the conditional probability that we really care about, the posterior $p(\theta | x_1, \dots, x_n)$; but you need the prior to get the posterior. As Leonard Jimmie Savage, a founder of Bayesian analysis, has famously said:

You can’t cook the Bayesian omelet
without breaking the Bayesian eggs.

Leonard Jimmy Savage

The ability to use prior information is a strength, especially when one has to make a decision based by very weak data. Later in the book we will see how priors can be used to convey the idea that




Figure 11: Prior for Bernoulli example.




Figure 12: Likelihood for Bernoulli model with $n = 10$ and $s = 9$.




Figure 13: Posterior for Bernoulli model with $n = 10$ and $s = 9$.

prior density

posterior density



Figure 14: Making a Bayesian omelet.

a functional relationship between two variables is in some sense smooth, and how this can prevent models from overfitting the data. Nevertheless, the subjective elements of a Bayesian analysis can complicate the reporting of scientific evidence, where objectivity is the ideal. One can argue that objectivity is simply unattainable, and that the supposedly objective alternatives to Bayesian learning just sweeps the subjective elements under the carpet. A more pragmatic Bayesian approach for scientific communication is presented in Section [Prior elicitation and objective Bayes](#) where priors are intentionally chosen to be ‘non-informative’, in the sense of having a minimal influence on the posterior.

There are also two aspects of a Bayesian approach that gives it a clear scientific character. The prior distribution is subjective, and therefore varies from person to person, but the rule that updates the beliefs with new data is objective: we *should* use Bayes’ theorem and the data *should* enter the updating *only through the likelihood function*. The word ‘should’ is emphasized here since one can mathematically derive this result from some simple axioms, and it can be proved to be the optimal way to process information; see [Bernardo and Smith \[2009\]](#) and Section [Bayesian learning and the likelihood principle](#). Second, one can prove that the effect of the prior vanishes asymptotically as the sample size n grows large; objectivity is attained by a **subjective consensus**: persons with wildly different priors will eventually reach the same posterior distribution as we collect more data. This result is given in chapter [Classification](#) and we will see an empirical demonstration of this effect already in the next chapter.

subjective consensus

EXERCISES

1. **Computer exercise.** This is the first problem.
2. This is the second problem.

Single-parameter models

Now that we know the basics of Bayesian updating of prior beliefs with new data, we can start to analyze models with a single parameter. This will allow to practice on deriving the posterior distribution in simple settings. The drawback of simple models is that they do not show anywhere near the full potential of Bayesian methods. But you need to crawl before you can walk, and some patience is required before we come to more useful models, such as regression and classification models in later chapters.

Bernoulli data

Let us return to iid Bernoulli data:

$$x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} \text{Bern}(\theta). \quad (10)$$

We first need a prior distribution $p(\theta)$ for θ . There are a number of ways to do **prior elicitation**, i.e. to extract a prior distribution from a person, for example an expert. Such methods involve ideas from psychology and usually consist of asking a series of questions to the expert, followed by checks for internal consistency of the elicited prior beliefs. One can in principle elicit any distribution, e.g. in the form of a histogram, but the most common approach is to first settle on a distributional family and then elicit the hyperparameters within the family. Since $\theta \in [0, 1]$, the **Beta distribution** is a suitable two-parameter family with quite a lot of flexibility; Figure 16 plots a few members of the Beta family. Note that $\text{Beta}(1, 1)$ is the **uniform distribution**. We will now show that the Beta family is particularly convenient as a prior for the iid Bernoulli model.

A nice feature of Bayesian inference is that one always know where to start. To derive the posterior distribution of a parameter θ we start with Bayes' theorem (9):

$$p(\theta | x_1, \dots, x_n) \propto p(x_1, \dots, x_n | \theta) p(\theta),$$

where $p(x_1, \dots, x_n | \theta) = \theta^s (1 - \theta)^f$ is the likelihood for iid Bernoulli

Beta distribution

$X \sim \text{Beta}(\alpha, \beta)$ for $X \in [0, 1]$.

$$p(x) = \frac{x^{\alpha-1} (1-x)^{\beta-1}}{B(\alpha, \beta)}$$

$$\mathbb{E}(X) = \frac{\alpha}{\alpha + \beta}$$

$$\mathbb{V}(X) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$$

$B(\alpha, \beta) = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha + \beta)}$, where $\Gamma(\alpha)$ is the Gamma function.

Figure 15: The beta distribution.




Figure 16: Some Beta distributions.

prior elicitation

Beta distribution

uniform distribution

Uniform distribution

$X \sim \text{Uniform}(a, b)$, $X \in [a, b]$.

$$p(x) = \frac{1}{b - a}$$

$$\mathbb{E}(X) = \frac{a + b}{2}$$

$$\mathbb{V}(X) = \frac{(b - a)^2}{12}$$

Figure 17: The uniform distribution.

data and $p(\theta)$ is the $\theta \sim \text{Beta}(\alpha, \beta)$ prior. So,

$$p(\theta|x_1, \dots, x_n) \propto \theta^s (1 - \theta)^f \frac{\theta^{\alpha-1} (1 - \theta)^{\beta-1}}{B(\alpha, \beta)} \quad (11)$$

$$\propto \theta^{\alpha+s-1} (1 - \theta)^{\beta+f-1}, \quad (12)$$

where the second line puts the Beta function $B(\alpha, \beta)$ into the missing proportionality constant. Note that $1/B(\alpha, \beta)$ is a multiplicative constant and *not* a function of θ and will therefore not affect the shape of the posterior distribution, just scale it vertically. In the final step will recover the normalizing constant so that $p(\theta|x_1, \dots, x_n)$ integrates to one over its support, as required. Now, from the pdf of the Beta distribution we see that the expression in (11) can be recognized as proportional to a Beta distribution. We see this as the expression is of the form $\theta^{a-1} (1 - \theta)^{b-1}$ where $a = \alpha + s$ and $b = \beta + f$. The posterior for θ is therefore the $\text{Beta}(\alpha + s, \beta + f)$ distribution and the missing proportionality constant in (11) is then known to be $1/B(\alpha + s, \beta + f)$. The prior-to-posterior updating for the Bernoulli model is summarized in Figure 18. Note that the random variables in the model are written with lowercase letters for simplicity.

Conjugate analysis - Bernoulli model

Model: $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} \text{Bern}(\theta)$

Prior: $\theta \sim \text{Beta}(\alpha, \beta)$

Posterior: $\theta | x_1, \dots, x_n \sim \text{Beta}(\alpha + s, \beta + f)$

where $s = \sum_{i=1}^n x_i$ and $f = n - s$.

Figure 18: Prior-to-Posterior updating for the Bernoulli data with a Beta prior.

Using a Beta prior for the Bernoulli parameter is convenient since the posterior distribution then belongs to the *same distributional family* as the prior distribution; the posterior is also a Beta distribution. The beta family is said to be *conjugate* to the Bernoulli model, or that the beta distribution is the **conjugate prior** for the Bernoulli model. Conjugate priors are easy to use since all we have to do when updating a Beta prior with Bernoulli data is to add the number of successes s to α and the number of failures f to β . The way that α and β enter the posterior also shows that the information in a $\text{Beta}(\alpha, \beta)$ prior corresponds to a prior dataset with α successes and β failures. We usually do not have an explicit prior sample at hand, and α and β need not even be integers, but we can nevertheless think about the prior information as being equivalent to an **imaginary prior sample**.

Similar conjugate results for several other models will be presented in this book, but there are many models for which a known conjugate prior do not exist. For such models, the posterior is often

conjugate prior

imaginary prior sample

not available in closed form, but several easy-to-use approximation or simulation methods are presented in later chapters.

It is interesting to compare a Bayesian analysis of Bernoulli data with the maximum likelihood estimator $\hat{\theta}_{MLE} = s/n$. A common **Bayes estimator**, or Bayesian point estimator, is the posterior mean $\mathbb{E}(\theta|x_1, \dots, x_n) = \frac{\alpha+s}{\alpha+\beta+n}$, which follows directly from the formula for the mean of a Beta distribution. Let us also assume a uniform prior for θ as some sort of non-informative prior, i.e. our prior is the Beta(1,1) distribution. Consider the case when we have observed no successes ($s = 0$) in a small number of trials n . We then have the quite unreasonable MLE of $\hat{\theta}_{MLE} = 0$, whereas the Bayes estimator is $\mathbb{E}(\theta|x_1, \dots, x_n) = 1/(n+2) > 0$. We will return to this example and the idea of a non-informative prior in Section [Prior elicitation and objective Bayes](#).

Bayes estimator




Figure 19: Bayesian analysis of $n = 10$ randomly chosen emails from the SpamBase data using three different priors. The likelihood is normalized.

Example: spam emails

The **SpamBase dataset** from the UCI repository¹ consists of 4601 emails that have been manually classified as *spam* (junk email) or *ham* (non-junk email). The dataset also contains a vector of covariates/features for each email, such as the number of capital letters or \$-signs; this information can be used to build a spam filter that automatically separates spam from ham. We will in this chapter only analyze the proportion of spam emails without using the covariates; we return to the more interesting case with features in the [Classification](#) chapter. So, let $x_1, \dots, x_n | \theta \stackrel{iid}{\sim} \text{Bern}(\theta)$ for the $n = 4601$ emails, where $x_i = 1$ if the email is spam and $x_i = 0$ for ham. The unknown quantity θ is the probability of spam.

To illustrate the incremental learning process in Bayesian learning we start off by analyzing only $n = 10$ randomly sampled emails, out of which $s = 4$ were spam. Figure 19 shows the posterior distribution of θ for three persons with very different priors. With only $n = 10$ data points, the three persons' posteriors are of course very different. The results in Figure 20 are based on $n = 100$ randomly sampled emails, including the 10 emails used in Figure 19. The posteriors

SpamBase dataset

¹ Dheeru Dua and Casey Graff. UCI machine learning repository, 2017. URL <http://archive.ics.uci.edu/ml/datasets/Spambase/>




Figure 20: Bayesian analysis of $n = 100$ randomly chosen emails from the SpamBase data using three different priors. The likelihood is normalized.

are now in rather close but not perfect agreement. Finally, Figure 21 shows the posterior for the full dataset with $n = 4601$; here there is a complete subjective consensus between the three persons that initially had very different beliefs about the spam probability.

From this dataset we have thus learned that around 40% of all emails are spam, and we are also quite certain about this percentage as the posterior distribution is very concentrated around 0.4. This information is not useful for building a spam filter where one instead needs the spam probability for each email to be a function of the text in that specific email (e.g. the number of \$-signs). We will achieve this in chapter [Classification](#) when we derive the posterior for a binary regression and use the methods in chapter [Prediction and Decision making](#) to construct Bayesian spam predictions from such a model.




Figure 21: Bayesian analysis of all $n = 4601$ emails from the SpamBase data using three different priors. The likelihood is normalized.

Bayesian learning and the likelihood principle

We will use the Bernoulli example to demonstrate an important feature of Bayesian learning. Consider the following three experiments, all resulting in s successes in n trials:

- **Experiment 1:** sample data from $X_1, \dots, X_n | \theta \sim \text{Bern}(\theta)$, where n is a predetermined number of trials.
Data: the outcome in each trial: x_1, \dots, x_n .
- **Experiment 2:** sample data from $X_1, \dots, X_n | \theta \sim \text{Bern}(\theta)$, where n is a predetermined number of trials.
Data: the total number of successes: $s = \sum_{i=1}^n x_i$

- **Experiment 3:** sample data from $X_i|\theta \sim \text{Bern}(\theta)$ until exactly s , a predetermined number of successes, have been obtained.

Data: the number of trials, n , until s successes have been obtained.

The above three experiments show that we need to be careful in defining exactly *which* data to use in the likelihood function. We know from before that the likelihood from Experiment 1 is

$$p(x_1, \dots, x_n|\theta) = \theta^s(1 - \theta)^{n-s}, \quad (13)$$

In the second experiment we only get to observe that there was s successes in n trials, but the exact sequence x_1, \dots, x_n is not recorded. So the data is here represented as the outcome of a random variable $S = \sum_{i=1}^n X_i \sim \text{Binom}(n, \theta)$. The likelihood for experiment 2 is therefore given by the binomial distribution

$$p(s) = \binom{n}{s} \theta^s (1 - \theta)^{n-s}. \quad (14)$$

This is different from the likelihood in Experiment 1 since the outcome $S = s$ can be obtained from several different observed data sequences x_1, \dots, x_n , each with exactly s successes. The exact number of such possible sequences is given by the binomial factor $\binom{n}{s}$.

Finally, the random variable in Experiment 3 is the number of performed trials, which follows the **negative binomial distribution**. The likelihood from Experiment 3 is therefore

negative binomial distribution

$$p(n) = \binom{n-1}{s-1} \theta^s (1 - \theta)^{n-s}. \quad (15)$$

The factor $\binom{n-1}{s-1}$ counts the number of ways we can order the $s - 1$ successes in the first $n - 1$ trials; we know that the n th trial must have been a success since the experiment terminated after n trials. Note that there are several versions of the negative binomial distribution depending on whether we count the number of trials or the number of failures until s successes.

Now, the likelihood functions in (13)-(15) differ only by a constant that does not depend on θ , i.e. the likelihoods are proportional. The likelihood for the j th experiment can therefore be written as $c_j f(\theta)$, where $f(\theta) = \theta^s (1 - \theta)^{n-s}$, $c_1 = 1$, $c_2 = \binom{n}{s}$ and $c_3 = \binom{n-1}{s-1}$. The posterior distribution of θ from the j th experiment is then by (8)

$$p_j(\theta|x_1, \dots, x_n) = \frac{c_j f(\theta) p(\theta)}{\int c_j f(\theta) p(\theta) d\theta} = \frac{f(\theta) p(\theta)}{\int f(\theta) p(\theta) d\theta}.$$

The posterior distribution for θ is therefore the same in all three experiments. It is now obvious that Bayesian inference always satisfies the following likelihood principle.

Definition. Likelihood principle. *Two experiments that result in (proportionally) equal likelihood functions should give the same inferences.*

Likelihood principle

Informally, the likelihood principle says that all relevant information in an experiment about θ is contained in the likelihood function. The importance of the likelihood principle is that it can be mathematically derived from two simpler principles that everyone hold as self evident. Hence the word *should* in the principle; see Casella and Berger [2002, ch. 6.2] for a discussion of this famous **Birnbaum's theorem**.

Birnbaum's theorem

Many frequentist methods violate the likelihood principle. The maximum likelihood *estimate* is easily seen to be $\hat{\theta}_{\text{MLE}} = s/n$ for all three experiments for a given data set. However, the sampling variability of the maximum likelihood *estimator*, $\mathbb{V}(\hat{\theta}_{\text{MLE}})$, will be different in Experiment 3 from that in Experiment 1 and 2. This is a consequence of the estimator being S/n in Experiment 1 and 2, but s/N in Experiment 3; note the difference in random variables (capital letters) in these estimators.

In summary, Bayesian inference *conditions on the observed data* and does not rely on repeated sampling properties. The data only enters through the likelihood function and Bayesian inference respects the likelihood principle.

Gaussian data - known variance

In this section we derive the posterior distribution for the mean in the iid Gaussian model $x_1, \dots, x_n | \theta, \sigma^2 \stackrel{\text{iid}}{\sim} N(\theta, \sigma^2)$. Since this chapter is about models with a single parameter we will assume the variance σ^2 to be known; this is rarely the case in practise and we return to the Gaussian model with both parameters unknown in Chapter [Multi-parameter models](#).

Uniform prior

We will first derive the posterior for a so called non-informative prior, i.e. a prior that is supposed to contain no, or at least very little, prior information. The most common non-informative prior for θ is a uniform distribution $p(\theta) = c$ for $\theta \in \mathbb{R}$ where $c > 0$ is a constant; the idea is that this distribution does not favor any particular value for θ . A uniform distribution over an unbounded space is not a proper distribution since $\int_{-\infty}^{\infty} p(\theta) d\theta = \infty$. It is nevertheless possible to use this somewhat strange prior since the resulting posterior is proper after observing a single data point. We can also think about the uniform prior as a limiting normal distribution with a variance that tends to infinity.

By Bayes' theorem, the posterior distribution for θ under a uniform prior is

$$\begin{aligned} p(\theta | x_1, \dots, x_n) &\propto p(x_1, \dots, x_n | \theta) p(\theta) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(x_i - \theta)^2\right) \cdot c \\ &\propto \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \theta)^2\right). \end{aligned}$$

Let $\bar{x}_n = \sum_{i=1}^n x_i$ be the sample mean, then

$$\sum_{i=1}^n (x_i - \theta)^2 = \sum_{i=1}^n (x_i - \bar{x} - (\theta - \bar{x}))^2 = \sum_{i=1}^n (x_i - \bar{x})^2 + n(\theta - \bar{x})^2,$$

since the cross term $2(\theta - \bar{x}) \sum_{i=1}^n (x_i - \bar{x}) = 0$. Note that the term $\sum_{i=1}^n (x_i - \bar{x})^2$ does not depend on θ and we therefore get

$$p(\theta | x_1, \dots, x_n) \propto \exp\left(-\frac{n}{2\sigma^2}(\theta - \bar{x})^2\right), \quad (16)$$

and hence that the posterior for θ can be recognized as

$$\theta | x_1, \dots, x_n \sim N\left(\bar{x}, \frac{\sigma^2}{n}\right).$$

Normal distribution

$$X \sim N(\mu, \sigma^2)$$

$$\text{Support: } X \in (-\infty, \infty)$$

$$p(x) = \frac{\exp\left(-\frac{1}{2\sigma^2}(x - \mu)^2\right)}{\sqrt{2\pi\sigma^2}}$$

$$\mathbb{E}(X) = \mu$$

$$\mathbb{V}(X) = \sigma^2$$

Figure 22: The Gaussian distribution. Confugate to Normal model.




Figure 23: Some Normal distributions.

Normal prior

Consider now a normal prior, $\theta \sim N(\mu_0, \tau_0^2)$; following Gelman et al. [2013] the subscript 0 is used to denote that these are **hyperparameters** in the prior, i.e. based on 0 observations. The user must decide the most probable value for θ , μ_0 , and also how sure she is by setting the prior standard deviation, τ_0 . One way to elicit these prior hyperparameters is to ask the user for a 95% probability interval for θ and then back out μ_0 and τ_0 ; see Exercise 2.

hyperparameters

By Bayes' theorem and the rewrite of the likelihood in (16) we have

$$p(\theta|x_1, \dots, x_n) \propto \exp\left(-\frac{n}{2\sigma^2}(\theta - \bar{x})^2\right) \times \exp\left(-\frac{1}{2\tau_0^2}(\theta - \mu_0)^2\right)$$

In Exercise 4 you are asked to complete the squares in this expression to prove that this expression is proportional to a normal density of the form given in Figure 24.

Conjugate analysis - Gaussian model with known variance

Model: $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} N(\theta, \sigma^2)$, σ^2 known

Prior: $\theta \sim N(\mu_0, \tau_0^2)$

Posterior: $\theta | x_1, \dots, x_n \sim N(\mu_n, \tau_n^2)$.

Posterior precision: $\frac{1}{\tau_n^2} = \frac{n}{\sigma^2} + \frac{1}{\tau_0^2}$

Posterior mean: $\mu_n = w\bar{x} + (1-w)\mu_0$, where $\bar{x} = \sum_{i=1}^n x_i$

Weight: $w = \frac{n/\sigma^2}{n/\sigma^2 + 1/\tau_0^2}$

Figure 24: Prior-to-Posterior updating for normal data with known variance and normal prior for the mean.

The normal prior is therefore conjugate to the normal model with known variance (i.e. a normal prior gives a normal posterior). The interpretation of the posterior mean μ_n and τ_n^2 in Figure 24 are quite intuitive. Note first that the expression for the posterior variance τ_n^2 is written in terms of precision = 1/variance. The first term $n/\sigma^2 = 1/(\sigma^2/n)$ is the precision in the data. This can be seen in several ways, for example by the sampling variance being $\mathbb{V}(\bar{x}) = \sigma^2/n$. Hence the formula for the posterior precision $\frac{1}{\tau_n^2} = \frac{n}{\sigma^2} + \frac{1}{\tau_0^2}$ can be read

Posterior precision = Data precision + Prior precision.

The posterior mean $\mu_n = w\bar{x} + (1-w)\mu_0$ is a weighted average of the data mean \bar{x} and the prior mean. The weight w on \bar{x} in Figure 24 is the data precision relative to the prior precision. The posterior therefore puts more emphasis on the data when n is large, σ small or τ_0 is large. It will not always be possible to get this clear a view of the prior-to-posterior updating in other models, but the same logic will apply also there.

Example: Internet connection speed

The maximum internet connection speed downstream in my home is 50 Mbit/sec. This maximum will typically never be reached, but my internet service provider (ISP) claims that the average speed is *at least* 20Mbit/sec. To test this, I collect a total of five measurements over the course of five consecutive using an speed testing internet service; I will call this the **Internet speed dataset**. The measurements are assumed to be $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} N(\theta, \sigma^2)$, where θ is the average speed; we ignore for simplicity that the measurements cannot be negative. The measurements are reported to have a standard deviation of $\sigma = 5$ by speed testing service. I will use a prior centered on the average claimed by the ISP, $\mu_0 = 20$, with a prior standard deviation of $\tau_0 = 5$. My prior beliefs are therefore that $\theta \in [10, 30]$ with approximately 95% probability.

Figure 25 displays the prior, normalized likelihood and posterior of θ based on only the first measurement $x_1 = 15.770$ Mbit/sec; the probability of interest $\Pr(\theta \geq 20 | x_1, \dots, x_n) \approx 0.275$ is marked out by the shaded orange region. Since the prior precision happenend to be equal to the data precision of a single observation, the weight on the data in the posterior mean μ_n is exactly $w = 0.5$. Figure 26 shows the results based on first two observations $x_1 = 15.770, x_2 = 20.501$. The posterior has shifted a little to the right and $\Pr(\theta \geq 20 | x_1, \dots, x_n)$ is now 0.333. Finally, using all $n = 5$ data points with $\bar{x} = 16.001$, Figure 27 shows that we are beginning to be rather confident that the ISP's claim that $\theta \geq 20$ is false since we now have $\Pr(\theta \geq 20 | x_1, \dots, x_n) \approx 0.051$. The weight w is now 0.833 so that data is starting to dominate the prior.

Figures 25-27 illustrate a situation where the posterior is computed by combining the prior at day 0, $N(\mu_0, \tau_0^2)$, with the likelihood for all x_1, \dots, x_n data points; hence the posterior on day n is computed as

$$p(\theta | x_1, \dots, x_n) \propto p(x_1, \dots, x_n | \theta) p(\theta). \quad (17)$$

We can however equally well compute this posterior by updating yesterday's posterior ($\theta | x_1, \dots, x_{n-1}$) with today's measurement x_n by

$$p(\theta | x_1, \dots, x_n) \propto p(x_n | \theta) p(\theta | x_1, \dots, x_{n-1}). \quad (18)$$

The updating in (17) and (18) give the same result, but (18) can be used sequentially in what is often called **online learning**, where "yesterday's posterior becomes today's prior". This online learning is illustrated in Figure 28 for the internet speed data.

The same online learning holds also for dependent data, e.g. time




Figure 25: Internet speed data. Posterior updating based on $n = 1$ observations. The orange shaded region marks out $\Pr(\theta \geq 20 | x_1, \dots, x_n)$




Figure 26: Internet speed data. Posterior updating based on $n = 2$ observations. The orange shaded region marks out $\Pr(\theta \geq 20 | x_1, \dots, x_n)$.




Figure 27: Internet speed data. Posterior updating based on $n = 5$ observations. The orange shaded region marks out $\Pr(\theta \geq 20 | x_1, \dots, x_n)$.




Figure 28: Internet speed data. Bayesian online learning.

online learning

series, as is easily proved as follows

$$\begin{aligned}
 p(\theta|x_1, \dots, x_n) &\propto p(x_1, \dots, x_n|\theta)p(\theta) \\
 &= p(x_n|\theta, x_1, \dots, x_{n-1})p(x_1, \dots, x_{n-1}|\theta)p(\theta) \\
 &\propto p(x_n|\theta, x_1, \dots, x_{n-1})p(\theta|x_1, \dots, x_{n-1}), \quad (19)
 \end{aligned}$$

where the second line follows from the decomposition results in Figure 29. For iid data we have the additional simplification $p(x_n|\theta, x_1, \dots, x_{n-1}) = p(x_n|\theta)$, hence showing the equivalence of (17) and (18).

By the same proof we also see that Bayesian methods are directly applicable in **batch learning**, where the posterior can be incrementally updated using batches of several observations, since for any $1 \leq m \leq n-1$

$$p(\theta|x_1, \dots, x_n) \propto p(x_{m+1}, \dots, x_n|\theta)p(\theta|x_1, \dots, x_m). \quad (20)$$

Implementing online or batch learning is straightforward for conjugate models since:

- any intermediate posterior $p(\theta|x_1, \dots, x_m)$ belongs to the same distribution family as the original prior $p(\theta)$ and
- the prior is conjugate to the likelihood for any data, and therefore also to the likelihood of the new batch $p(x_{m+1}, \dots, x_n|\theta)$.

In the case of the iid normal model with known variance we have the recursions for observation $i = 1, 2, \dots$

$$\begin{aligned}
 \frac{1}{\tau_i^2} &= \frac{1}{\sigma^2} + \frac{1}{\tau_{i-1}^2} \\
 w_i &= \frac{\sigma^{-2}}{\sigma^{-2} + \tau_{i-1}^{-2}} \\
 \mu_i &= w_i x_i + (1 - w_i) \mu_{i-1}.
 \end{aligned}$$

When the prior is not conjugate one has to resort to numerical methods that can be more or less computationally attractive in online mode; see in the chapters [Posterior simulation](#) and [Variational inference](#).

Poisson data

Count data $X \in 0, 1, 2, \dots$ is a quite frequently occurring data type in many applications; some examples are the number of software bugs, the number of lethal car accidents in a region, or the number of scooters available at a given pick-up station. The most commonly used model for count data is the **Poisson distribution**. The mean and variance of a Poisson variable are always equal, which can be

batch learning

Decomposing distributions

For two random variables X, Y

$$p(x, y) = p(y|x)p(x)$$

For n random variables

$$p(x_1, \dots, x_n) = p(x_1)p(x_2|x_1) \times \dots \times p(x_n|x_1, \dots, x_{n-1})$$

and conditional on θ

$$p(x_1, \dots, x_n|\theta) = p(x_1|\theta) \times \dots \times p(x_n|x_1, \dots, x_{n-1}, \theta)$$

Figure 29: Marginal-Conditional decomposition of a joint distribution.

Poisson distribution

$X \sim \text{Pois}(\theta)$ for $X \in 0, 1, 2, \dots$

$$p(x) = \frac{\theta^x e^{-\theta}}{x!}$$

$$\mathbb{E}(X) = \theta$$

$$\mathbb{V}(X) = \theta$$

Figure 30: The Poisson distribution.




Figure 31: Some Poisson distributions.

Poisson distribution

restrictive in some application, but the model often fits many real datasets surprisingly well or can be extended to do so.

Figure 32.

The likelihood function for $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} \text{Pois}(\theta)$, is

$$p(x_1, \dots, x_n | \theta) = \prod_{i=1}^n \frac{\theta^{x_i} e^{-\theta}}{x_i!} \propto \theta^{\sum_{i=1}^n x_i} e^{-n\theta}. \quad (21)$$

Comparing the functional form of the likelihood in (21) with a list of common probability distributions we can see that the likelihood from iid Poisson data looks very much like a **Gamma distribution** in θ . Even more, the form of the Gamma distribution tells us that a Gamma prior may indeed combine nicely with this likelihood. So let us try if $\theta \sim \text{Gamma}(\alpha, \beta)$ is conjugate to the iid Poisson model:

$$\begin{aligned} p(\theta | x_1, \dots, x_n) &\propto p(x_1, \dots, x_n | \theta) p(\theta) \\ &\propto \theta^{\sum_{i=1}^n x_i} e^{-n\theta} \cdot \theta^{\alpha-1} e^{-\beta\theta} \\ &= \theta^{\alpha + \sum_{i=1}^n x_i - 1} e^{-(\beta+n)\theta}, \end{aligned}$$

where we have directly written up the $\text{Gamma}(\alpha, \beta)$ prior without normalization constant. This expression is indeed proportional to a Gamma distribution and we have the following result:

Conjugate analysis - Poisson model

Model: $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} \text{Pois}(\theta)$

Prior: $\theta \sim \text{Gamma}(\alpha, \beta)$

Posterior: $\theta | x_1, \dots, x_n \sim \text{Gamma}(\alpha + \sum_{i=1}^n x_i, \beta + n)$

Example: internet auction data

The **eBayCoin dataset** collected by Wegmann and Villani [2011] and made available in the UCI repository² consist of data from 1000 eBay auctions of collectors coins. For each auction, the dataset records the final price of the auctioned coin, the number of bidder in the auction and a number of covariates such as the quality of the sold coin, the lowest price that the seller would agree to sell for etc. We will here analyze the number of bidders using an iid Poisson model without covariates. We return to this dataset in Chapter Classification where we make use of the covariates in a Poisson regression model for predicting the number of bidders.

To compute the posterior distribution for θ , the average number of bidders in an auction we need the summary statistic $\sum_{i=1}^n x_i = 3635$.

Gamma distribution

$X \sim \text{Gamma}(\alpha, \beta)$ for $X > 0$.

$$p(x) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}$$

$$\mathbb{E}(X) = \frac{\alpha}{\beta}$$

$$\mathbb{V}(X) = \frac{\alpha}{\beta^2}$$

Figure 32: Gamma distribution.

Gamma distribution




Figure 33: Some Gamma distributions.

Figure 34: Prior-to-Posterior updating for the Poisson data with a Gamma prior.

eBayCoin dataset

²<http://archive.ics.uci.edu/ml/datasets/eBayCoin/>

The sample mean in the $n = 1000$ auctions is therefore $\bar{x} = 3.635$ bidders per auction. I will use the gamma prior with $\alpha = 2$ and $\beta = 1/2$ since this implies a prior mean of $\mathbb{E}(\theta) = 4$ and prior standard deviation of $S(\theta) = 2.283$, which I find matches quite well with my prior beliefs. This prior and the posterior updated with data from $n = 1000$ auctions are shown in Figure 35. Note the different scales on the horizontal axis. We are now more or less certain that the average number of bidders is in the interval $\theta \in [3.4, 3.9]$.




Figure 35: Bayesian analysis of the numbers of bidders in $n = 1000$ eBay coin auctions.

Figure 36 a) plots the fitted Poisson distribution with θ set equal to the posterior mean against the observed data. It is obvious that the Poisson distribution is too restrictive as the fit is terrible.




Figure 36: Assessing the fit of the Poisson model with the posterior mean estimate of θ .

The poor fit can be attributed to the heterogeneity of the auctions. For example, some of the auctions had a high so called reservation price, i.e. the lowest price that the seller is willing sell for, while other auctions had a very low reservation price. It is expected that a high reservation price discourages bidders from entering the auction.

To explore the effect of the reservation price we split the data into low and high reservation price auctions, and analyze the two auction types separately. The prior-to-posterior updating is shown in Figure 37; the priors now reflect that θ is likely to be larger for the auctions with low reservation prices. The posteriors are clearly different in the two subpopulations. The Poisson model fits better on the two subpopulations as shown in Figure 36 b) and c), but it is not perfect. We will return to this dataset in Chapter Regression using a Poisson regression with the reservation price as covariate as well as other auction specific covariates.




Figure 37: eBay auctions. Bayesian analysis of the numbers of bidders in $n = 550$ auctions with a low reservation price and $n = 450$ auctions with a high reservation price.

Summarizing a posterior distribution

The posterior distribution for models with a single parameter are easily plotted and gives a complete visual quantification of uncertainty. Starting from the next chapter, our models will typically contain more than one parameter, and not seldom quite many. It is then impractical to plot the whole posterior distribution and we will now explore some commonly used numerical summaries of the posterior, for example a point estimate and posterior probability intervals.

A point estimate of θ summarizes the posterior with a single point. The three most commonly used Bayesian point estimates are:

- The posterior mean $\hat{\theta}_{\text{mean}} \equiv \mathbb{E}(\theta|x_1, \dots, x_n)$.
- The posterior median $\hat{\theta}_{\text{med}}$, i.e. the 50th quantile of $p(\theta|x_1, \dots, x_n)$.
- The posterior mode $\hat{\theta}_{\text{mode}} \equiv \arg \max_{\theta \in \Theta} p(\theta|x_1, \dots, x_n)$.

We will see in chapter [Prediction and Decision making](#) that the choice of point estimate can be formalized as a decision problem.

A point estimate says nothing about the variability in the posterior. One way to quantify the uncertainty is the posterior standard deviation $S(\theta|x_1, \dots, x_n) = \sqrt{\mathbb{V}(\theta|x_1, \dots, x_n)}$.

EXAMPLE: POISSON DATA. As we saw earlier the posterior for the mean θ of a Poisson distribution with a $\theta \sim \text{Gamma}(\alpha, \beta)$ prior is $\theta|x_1, \dots, x_n \sim \text{Gamma}(\alpha + \sum_{i=1}^n x_i, \beta + n)$. From properties of the Gamma distribution, the posterior mean estimate is hence $(\alpha + \sum_{i=1}^n x_i) / (\beta + n)$ and the posterior variance is $(\alpha + \sum_{i=1}^n x_i) / (\beta + n)^2$. For the eBay data [Example: internet auction data](#) we have $\mathbb{E}(\theta|x_1, \dots, x_n) = \frac{2+3635}{0.5+1000} \approx 3.635$ bidders and $S(\theta|x_1, \dots, x_n) = \sqrt{\frac{2+3635}{(0.5+1000)^2}} \approx 0.060$.

Before presenting how to summarize a posterior by an interval, let us first informally recall the definition of a frequentist confidence interval. A 95% *confidence interval* for a parameter θ is a random interval $[l(X_1, \dots, X_n), u(X_1, \dots, X_n)]$ that contains the true θ in 95% of all possible datasets X_1, \dots, X_n from the data generating process.




Figure 38: Three common point estimates for summarizing a posterior.

As usual with frequentist methods we are guaranteed a long run performance over all possible datasets, but the realized interval $[l(x_1, \dots, x_n), u(x_1, \dots, x_n)]$ either does or does not cover the true θ .

A Bayesian interval is defined in a much more direct way, and is conditional on the actually observed dataset. This simpler definition is possible since the posterior is a probability distribution; we have broken the Bayesian eggs and can enjoy the omelet. A 95% posterior **credibility interval** for $\theta \in \Theta \subset \mathbb{R}$ is an interval $[l, u] \subset \Theta$ such that $\Pr(\theta \in [l, u] | x_1, \dots, x_n) = 0.95$, i.e. an interval that contains 95% of the posterior probability mass. We can generalize this to a more general region than an interval, for example a union of disjoint intervals, and of course to other probability coverages than 95%.

credibility interval

There are many ways to construct an credibility interval with a certain coverage probability. An **equal tail credibility interval** is an interval that cuts off equal probability in the left and right tail; for example, a 95% interval sets l and u to the 2.5% and 97.5% posterior quantile, respectively. Another popular interval construction is the highest posterior density (HPD) region which, as the name suggest, is made up of the θ values with highest posterior density. We use the word *region* instead of interval here since HPD regions need not be intervals. Here is the definition.

equal tail credibility interval

Definition (HPD region). A *Highest Posterior Density (HPD) region* for $\theta \in \Theta$ with coverage probability γ is a region $R \subset \Theta$ such that:

Highest Posterior Density (HPD) region

- $\Pr(\theta \in R | x_1, \dots, x_n) = \gamma$ and
- $p(\theta_{\text{in}} | x_1, \dots, x_n) \geq p(\theta_{\text{out}} | x_1, \dots, x_n)$ for all $\theta_{\text{in}} \in R$ and $\theta_{\text{out}} \notin R$.

Figure 39 illustrates the difference between equal tail intervals and HPD regions for some example densities. Note how the equal tail interval construction can exclude θ values that actually have highest posterior density (middle graph) and how HPD regions can be disconnected (righthand graph).




Figure 39: HPD regions vs equal tail intervals.

A disadvantage of HPD regions is that they are not invariant to reparametrization; the HPD for a transformed parameter $\eta(\theta)$ is

typically not a direct map of the HPD for θ .

EXAMPLE: POISSON DATA The 95% equal tail interval for the mean number of bidders is $[3.518, 3.754]$ which is virtually indistinguishable from the HPD interval $[3.517, 3.754]$ since the posterior is essentially symmetric, see Figure 40.

Exponential Family and Sufficiency*

This section presents the concept of sufficient statistics and the exponential family of distributions, with particular emphasis on their role in Bayesian learning. While these concepts are very important in statistics, this starred section can be skipped at first reading, but should be read before the generalized linear models in Chapter Classification, where the exponential family plays a prominent role.

Sufficient statistics

In all models covered so far in this book, the dataset, (x_1, \dots, x_n) , has only entered the likelihood through some low-dimensional summary statistic; for example the number of successes $s = \sum_{i=1}^n x_i$ in the Bernoulli model, the sample mean $\bar{x} = n^{-1} \sum_{i=1}^n x_i$ in the Gaussian model, and the sum of counts, $\sum_{i=1}^n x_i$, in the Poisson model. Note that we did not choose this data reduction, it just turned out that the likelihood only depended on the summarizing statistic; the statistic captured all the relevant information in the sample. In all of the above examples, the statistic was one-dimensional. In other models more than a single dimension is needed to compress the dataset, and we let the vector-valued function $\mathbf{t}(x_1, \dots, x_n) \rightarrow \mathbb{R}^k$ denote the statistic in general, where k is the dimension of reduction.

The following definition captures the idea that a statistic may contain *all* relevant information in the data about a parameter θ .

Definition. Sufficient statistic. A statistic $\mathbf{t}(X_1, \dots, X_n)$ is sufficient for θ if the conditional distribution of the sample X_1, \dots, X_n given the value of the statistic $\mathbf{t}(X_1, \dots, X_n)$ does not depend on θ .

The sufficiency of a statistic can be checked by the following lemma; see Casella and Berger [2002] for a proof.

Lemma 1. Factorization criterion. A statistic $\mathbf{t}(x_1, \dots, x_n)$ is sufficient for a parameter θ if and only if the likelihood can be factorized as

$$p(x_1, \dots, x_n | \theta) = h(x_1, \dots, x_n) f(\mathbf{t}(x_1, \dots, x_n); \theta), \quad (22)$$

where $h(x_1, \dots, x_n)$ does not depend on θ and $f(\mathbf{t}; \theta)$ is a function of the data only through the sufficient statistic $\mathbf{t}(x_1, \dots, x_n)$.




Figure 40: 95% credibility intervals for the Gamma posterior in the eBay auction data.

Sufficient statistic

Factorization criterion

The idea behind sufficient statistics is so appealing that it is often formulated as a desired inference principle similar to the likelihood principle presented in the section [Bayesian learning and the likelihood principle](#).

Definition. Sufficiency principle. If $\mathbf{t}(X_1, \dots, X_n)$ is a sufficient statistic for θ then any inference about θ should depend on the sample x_1, \dots, x_n only through the value $\mathbf{t}(x_1, \dots, x_n)$.

Sufficiency principle

Theorem 1. Bayesian learning satisfies the sufficiency principle.

Proof. If $\mathbf{t}(x_1, \dots, x_n)$ is a sufficient statistic for θ then by Lemma 1

$$\begin{aligned} p(\theta|x_1, \dots, x_n) &= \frac{p(x_1, \dots, x_n|\theta)p(\theta)}{\int p(x_1, \dots, x_n|\theta)p(\theta)d\theta} \\ &= \frac{h(x_1, \dots, x_n)f(\mathbf{t}(x_1, \dots, x_n); \theta)p(\theta)}{\int h(x_1, \dots, x_n)f(\mathbf{t}(x_1, \dots, x_n); \theta)p(\theta)d\theta} \\ &= \frac{f(\mathbf{t}(x_1, \dots, x_n); \theta)p(\theta)}{\int f(\mathbf{t}(x_1, \dots, x_n); \theta)p(\theta)d\theta}, \end{aligned}$$

which only depends on the data through the sufficient statistic $\mathbf{t}(x_1, \dots, x_n)$. □

Exponential family

All models considered so far are part of the large and important exponential family of distributions. A random variable X follows a distribution in the (one-parameter) **exponential family** if its density can be written in the form

exponential family

$$p(x|\theta) = h(x) \exp(\eta(\theta)t(x) - A(\theta)), \text{ for } x \in \mathcal{X}, \quad (23)$$

where $h(x)$ is a function of only x and $A(\theta)$ is a function of only θ . The support \mathcal{X} is not allowed to depend on θ , so that for example the $\text{Uniform}(0, \theta)$ distribution does not belong to the exponential family. The function $\eta(\theta)$ is called the **natural parameter** and is an invertible transformation of the parameter θ . Here are some examples.

natural parameter

EXAMPLE: POISSON DISTRIBUTION. The $\text{Pois}(\theta)$ distribution can be rewritten as follows

$$p(x|\theta) = \frac{\theta^x e^{-\theta}}{x!} = \frac{e^{x \ln \theta} e^{-\theta}}{x!} = \frac{1}{x!} \exp(x \ln \theta - \theta),$$

which is in the exponential family with $h(x) = (x!)^{-1}$, $A(\theta) = \theta$, $\eta(\theta) = \ln \theta$ and $t(x) = x$. Note in particular that the natural parameter is the logarithm of the Poisson mean, $\eta(\theta) = \ln \theta$.

EXAMPLE: BERNOULLI DISTRIBUTION. The $\text{Bern}(\theta)$ distribution can also be written as an exponential family:

$$p(x|\theta) = \theta^x (1-\theta)^{1-x} = \left(\frac{\theta}{1-\theta}\right)^x (1-\theta) = \exp\left(\eta(\theta)x - A(\theta)\right),$$

where $\eta(\theta) = \ln\left(\frac{\theta}{1-\theta}\right)$, $A(\theta) = \ln\left(\frac{1}{1-\theta}\right)$, $t(x) = x$ and $h(x) = 1$. The natural parameter for the Bernoulli distribution is therefore the log-odds, $\ln\left(\frac{\theta}{1-\theta}\right)$.

The normal distribution and many other distributions can similarly be shown to belong to the exponential family; but not all do, for example the **student- t distribution**. We will use $\text{ExpFam}(\theta)$ as a generic notation for a distribution in the exponential family, leaving the specific $h(x)$, $A(\theta)$, $\eta(\theta)$ and $t(x)$ functions implicit.

The likelihood function for iid data from an $\text{ExpFam}(\theta)$ distribution is

$$p(x_1, \dots, x_n|\theta) = \left[\prod_{i=1}^n h(x_i)\right] \exp\left(\eta(\theta) \sum_{i=1}^n t(x_i) - nA(\theta)\right). \quad (24)$$

Lemma 1 can be directly used to show that $\sum_{i=1}^n t(x_i)$ is a sufficient statistic for θ . In the next chapter we will see a multiparameter version of the exponential family with a vector of k sufficient statistics. The Pitman–Koopman–Darmois theorem [Bernardo and Smith, 2009] proves that among distributions whose support does not depend on θ , only the exponential family have sufficient statistics of fixed dimension, i.e. the dimension k does not depend on the size of the data, n (or at least is bounded).

The exponential family has several other attractive properties [Sundberg, 2019]. One property of particular interest here is that a conjugate prior always exists for models in the exponential family. In fact, the following family of priors is conjugate to the exponential family likelihood in (24)

$$p(\theta) = H(\tau_0, \nu_0) \exp\left(\eta(\theta)\tau_0 - \nu_0 A(\theta)\right), \quad (25)$$

where $H(\tau_0, \nu_0)$ is the normalizing constant. Note that this prior has two hyperparameter τ_0 and ν_0 that needs to be set by the user. We will use the symbol $\theta \sim \text{ExpFamConj}(\tau_0, \nu_0)$ for this prior distribution, where it must be remembered that the form of the prior depends on which specific exponential family member the prior is conjugate to, i.e. it depends on $\eta(\theta)$ and $A(\theta)$.

EXAMPLE: BERNOULLI MODEL. It was shown above that $\eta(\theta) = \ln\left(\frac{\theta}{1-\theta}\right)$ and $A(\theta) = \ln\left(\frac{1}{1-\theta}\right)$, for Bernoulli data. The prior in (25) is

student- t distribution

Student- t distribution

$X \sim t(\mu, \sigma, \nu)$ for $X \in (-\infty, \infty)$

$$p(x) = \frac{\Gamma((\nu+1)/2)}{\Gamma(\nu/2)\sqrt{\pi\nu\sigma^2}} \times \left(1 + \frac{1}{\nu} \left(\frac{x-\mu}{\sigma}\right)^2\right)^{-\frac{\nu+1}{2}}$$

$$\mathbb{E}(X) = \mu \text{ if } \nu > 1$$

$$\mathbb{V}(X) = \sigma^2 \frac{\nu}{\nu-2} \text{ if } \nu > 2$$

Figure 41: The student- t distributions.




Figure 42: Some Student- t distributions.

therefore

$$\begin{aligned} p(\theta) &\propto \exp\left(\eta(\theta)\tau_0 - \nu_0 A(\theta)\right) \\ &= \exp\left(\ln\left(\frac{\theta}{1-\theta}\right)\tau_0 - \nu_0 \ln\left(\frac{1}{1-\theta}\right)\right) \\ &\propto \theta^{\tau_0} (1-\theta)^{\nu_0 - \tau_0}, \end{aligned}$$

which is proportional to the $\text{Beta}(\tau_0, \nu_0 - \tau_0)$ distribution. The parametrization in (25) is hence interpreted as the information from a (imaginary) prior sample of τ_0 success in ν_0 trials. The $\text{Beta}(\alpha, \beta)$ prior from before expresses instead the prior information as a sample of α success and β failures.

The posterior distribution for θ in the exponential family with a conjugate prior is obtained by multiplying the likelihood in (24) with prior (25)

$$p(\theta|x_1, \dots, x_n) \propto \exp\left[\eta(\theta)\left(\tau_0 + \sum_{i=1}^n t(x_i)\right) - (\nu_0 + n)A(\theta)\right],$$

which is of the form ExpFamConj , but with updated hyperparameters: $\tau_0 \Rightarrow \tau_0 + \sum_{i=1}^n t(x_i)$ and $\nu_0 \Rightarrow \nu_0 + n$. We summarize this in Figure 43.

Conjugate analysis from iid exponential family data

Model: $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} \text{ExpFam}(\theta)$

Prior: $\theta \sim \text{ExpFamConj}(\tau_0, \nu_0)$

Posterior: $\theta | x_1, \dots, x_n \sim \text{ExpFamConj}(\tau_0 + \sum_{i=1}^n t(x_i), \nu_0 + n)$

This result shows that we can think quite generally about ν_0 as the (imaginary) prior sample size and τ_0 as the prior data compressed by the sufficient statistic. For example, in the Poisson model the information in the conjugate prior equals a prior sample of ν_0 data points with a mean count of τ_0 / ν_0 .

EXERCISES

1. Let $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} \text{Expon}(\theta)$ be exponentially distributed data. Show that the Gamma distribution is the conjugate prior for this model.
2. I determined my normal prior in the internet speed data example by specifying the prior mean θ_0 and standard deviation τ_0 . Assume that another person instead specified a 95% prior probability interval for θ as $[20, 30]$. Use this information to determine that person's normal prior, i.e. compute θ_0 and τ_0 for this person.

Figure 43: Prior-to-Posterior updating for iid exponential family data with a conjugate prior.

Exponential distribution

$X \sim \text{Expon}(\theta)$ for $X \in (0, \infty)$

$$p(x) = \theta e^{-\theta x}$$

$$\mathbb{E}(X) = 1/\theta$$

$$\mathbb{V}(X) = 1/\theta^2$$

Figure 44: The exponential distribution.




Figure 45: Some Exponential distributions.

3. (a) Let x_1, \dots, x_{10} be a sample with $\bar{x} = 1.873$. Assume the model $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} N(\theta, 1)$ and the prior $\theta \sim N(0, 5)$. Compute the posterior distribution of θ .
- (b) You now get hold of a second sample $y_1, \dots, y_{10} | \theta \stackrel{\text{iid}}{\sim} N(\theta, 2)$, where θ is the same quantity as in (a) but the measurements have a larger variance. The sample mean in this second sample is $\bar{y} = 0.582$. Compute the posterior distribution of θ using both samples (the x 's and the y 's) under the assumption that the two samples are independent.
- (c) You finally obtain a third sample $z_1, \dots, z_{10} | \theta \stackrel{\text{iid}}{\sim} N(\theta, 3)$, with mean $\bar{z} = 1.221$. Unfortunately, the measuring device for this latter sample was defective and any measurement above 3 was recorded as exactly 3. There were two such measurements. Give an expression for the unnormalized posterior distribution (likelihood \times prior) for θ based on all three samples (x, y and z). If you have computer available you may plot this unnormalized posterior over a grid of θ values. *Hint: the posterior distribution is not normal anymore when the measurements are truncated at 3.*
4. Derive the posterior distribution for the normal model with a normal prior in Figure 24. *Hint: complete the square.*
5. (a) Let $x_1, \dots, x_n | \theta \sim \text{Uniform}(\theta - 1/2, \theta + 1/2)$. Let $\hat{\theta} = \bar{x}$ be an estimator of θ . Derive an expression for the sampling variance of $\hat{\theta}$.
- (b) Derive the posterior distribution for θ assuming a uniform prior distribution. *Hint: once you have observed some data, some values for θ are no longer possible.*
- (c) Assume that you have observed three data observations: $x_1 = 1.1, x_2 = 2.09, x_3 = 1.4$. What would a frequentist conclude about θ ? What would a Bayesian conclude? Discuss.
6. Show that the $N(\mu, 1)$ distribution belongs to the exponential family.

NOTEBOOK EXERCISES

1. See the notebook [single-parameter models](#).

Multi-parameter models

Joint posterior distributions

Most models have more than one parameter, and many models are incredibly rich on parameters. Datasets are increasingly rapidly in size and makes it possible to estimate increasingly more complex models. To explore how Bayesian methods can be used in multi-parameter models we first return in this chapter to the iid $N(\theta, \sigma^2)$, but now in the more realistic setting where both θ and σ^2 are unknown parameters. In later chapters we will tackle regression and classification models where each covariate (input) x_k affects the response (output) y through a regression coefficient β_k ; hence in a regression with K covariates we have K regression coefficients β_1, \dots, β_K .

Consider a general probability model $p(x_1, \dots, x_n | \theta_1, \dots, \theta_K)$ with K parameters for a dataset x_1, \dots, x_n ; for example the iid normal model where $\theta_1 = \mu$ and $\theta_2 = \sigma^2$. Bayesian learning proceeds exactly as with a single parameter, except that the prior and posterior distribution are now both multidimensional joint distributions. Figure 46 gives an illustration of a bivariate ($K = 2$) normal distribution.

Using Bayes' theorem in proportional form, the **joint posterior distribution** $p(\theta_1, \dots, \theta_K | x_1, \dots, x_n)$ is given by

$$p(\theta_1, \dots, \theta_K | x_1, \dots, x_n) \propto p(x_1, \dots, x_n | \theta_1, \dots, \theta_K) p(\theta_1, \dots, \theta_K),$$

where $p(\theta_1, \dots, \theta_K)$ is a multidimensional prior distribution and $p(x_1, \dots, x_n | \theta_1, \dots, \theta_K)$ is the likelihood function; Note that the likelihood function is now a **likelihood surface** in the sense that it is a function of several parameters, $\theta_1, \dots, \theta_K$.

To keep the notation simpler we often use vector notation and write $\boldsymbol{\theta} \equiv (\theta_1, \dots, \theta_n)$ and $\mathbf{x} \equiv (x_1, \dots, x_n)$. The multivariate Bayes' theorem can then be expressed as

$$p(\boldsymbol{\theta} | \mathbf{x}) \propto p(\mathbf{x} | \boldsymbol{\theta}) p(\boldsymbol{\theta}). \quad (26)$$




Figure 46: Surface and contour plot of the bivariate normal distribution. The contour levels contain 25, 50, 75, 95 and 99% of the probability mass, respectively.

joint posterior distribution

likelihood surface

Marginalization

The joint posterior distribution $p(\theta|\mathbf{x})$ contains all posterior information about θ , but is obviously hard to visualize in the same way as we did for single-parameter models. In many cases we are also most interested in a subset of parameters, and the other parameters are only needed to model the data well but are of no real interest. Such parameters are just a nuisance when presenting inferences and are therefore often called **nuisance parameters**. Getting rid of nuisance parameters is very difficult in a non-Bayesian setting, for example when using maximum likelihood estimation. So what is the Bayesian solution to this dilemma?

Nuisance parameters can be handled in a very natural way in a Bayesian approach since the posterior distribution is a probability distribution for θ . We can therefore just integrate out, or marginalize out, the nuisance parameters just as in ordinary probability calculus. Take a simple example where $\theta = (\theta_1, \theta_2)$ and assume that the parameter of interest is θ_1 whereas θ_2 is considered a nuisance parameter; θ_1 could for example be the mean of iid Gaussian model and θ_2 the variance. The marginal posterior of θ_1 is then

$$p(\theta_1) = \int p(\theta_1, \theta_2) d\theta_2,$$

where the integration is over the full support of θ_2 . Figure 47 illustrates the marginalization concept. Using the decomposition $p(\theta_1, \theta_2) = p(\theta_1|\theta_2)p(\theta_2)$ we can alternatively express this as

$$p(\theta_1) = \int p(\theta_1|\theta_2)p(\theta_2)d\theta_2,$$

which shows that marginalization is achieved by averaging over the values of θ_2 with weights given by $p(\theta_2)$.

More generally, with more than two parameters, partition the elements of θ into two vectors, θ_a and θ_b . The marginal posterior of θ_a is the obtained by marginalizing out θ_b from the joint posterior

$$p(\theta_a) = \int \cdots \int p(\theta_a, \theta_b) d\theta_b. \quad (27)$$

We will see examples of marginalization in the following sections.

Gaussian data with unknown variance

The previous chapter analyzed iid normal data $x_1, \dots, x_n | \theta, \sigma^2 \stackrel{\text{iid}}{\sim} N(\theta, \sigma^2)$ under the usually unrealistic assumption that σ^2 is known. Let us now tackle the case where both parameters are unknown. It

nuisance parameters




Figure 47: Contour plot of the bivariate normal distribution in Figure 46 along with the marginal distributions.

turns out that the conjugate prior for this model has dependence between θ and σ , so we will describe the prior using the decomposition $p(\theta|\sigma^2)p(\sigma^2)$ as follows

$$\theta|\sigma^2 \sim N(\mu_0, \sigma^2/\kappa_0) \quad (28)$$

$$\sigma^2 \sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2). \quad (29)$$

The marginal conjugate prior for σ^2 involves a new distribution, the **scaled inverse chi-squared distribution**, denoted by $\text{Inv-}\chi^2(\nu_0, \sigma_0^2)$; see Figure 48. This distribution is a specific parametrization of the **inverse Gamma distribution**. The name comes from the characterization

$$X \sim \chi_\nu \Rightarrow Y = \nu\tau^2 \frac{1}{X} \sim \text{Inv-}\chi^2(\nu, \tau^2),$$

so that a $\text{Inv-}\chi^2(\nu, \tau^2)$ variable really is an inverted χ_ν^2 variable scaled by $\nu\tau^2$. Note that the parameter τ^2 is close to the mean when ν is large. The mode is $\nu\tau^2/(\nu+2)$, so τ^2 is somewhere between the mode and the mean. We will therefore call τ^2 the location of $\text{Inv-}\chi^2(\nu, \tau^2)$, or sometimes just sloppily as "our best guess".

The conjugate prior in (28) is specified via the four prior hyperparameters:

- μ_0 - the prior mean for θ
- κ_0 - the number of prior data observations for θ
- σ_0^2 - the prior location of σ^2
- ν_0 - the prior degrees of freedom for σ^2 .

Note that, similar to the conjugate prior for the exponential family, we are only *interpreting* κ_0 as the number of prior observations. The prior may not actually be based on previous data, but the information in the prior $\theta|\sigma^2 \sim N(\mu_0, \sigma^2/\kappa_0)$ has the equivalent strength of an imaginary prior sample of κ_0 observations from a data generating process with variance σ^2 .

Figure 49 shows that the posterior is indeed in the same form as the prior in (28), as required for a conjugate prior. There is a lot of greek letters in Figure 49, but note that the same sort of intuition applies here as in the case with a known variance in Chapter [Single-parameter models](#):

- the posterior mean μ_n is a weighted average of the data mean \bar{x} and the prior mean μ_0
- the weight on the data $w = n/(\kappa_0 + n)$ is close to one when either the data is informative (large n) or the prior is weak (small κ_0)

scaled inverse chi-squared distribution

inverse Gamma distribution

Inv- χ^2 distribution

$X \sim \text{Inv-}\chi^2(\nu, \tau^2), X \in (0, \infty)$

$$p(x) = \frac{(\tau^2\nu/2)^{\nu/2}}{\Gamma(\nu/2)} \frac{\exp\left(-\frac{\nu\tau^2}{2x}\right)}{x^{1+\nu/2}}$$

$$\mathbb{E}(X) = \frac{\nu}{\nu-2} \tau^2$$

$$\mathbb{V}(X) = \frac{2\nu^2\tau^4}{(\nu-2)^2(\nu-4)}$$




Figure 48: Some Scaled-Inv-Gamma distributions.

Gaussian iid data with conjugate prior**Model:** $x_1, \dots, x_n | \theta, \sigma^2 \stackrel{\text{iid}}{\sim} N(\theta, \sigma^2)$ **Prior:** $\theta | \sigma^2 \sim N(\mu_0, \sigma^2 / \kappa_0)$
 $\sigma^2 \sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2)$ **Posterior:** $\theta | \sigma^2, x_1, \dots, x_n \sim N(\mu_n, \sigma^2 / \kappa_n)$
 $\sigma^2 | x_1, \dots, x_n \sim \text{Inv-}\chi^2(\nu_n, \sigma_n^2)$

$$\mu_n = w\bar{x} + (1 - w)\mu_0$$

$$w = \frac{n}{\kappa_0 + n}$$

$$\kappa_n = \kappa_0 + n$$

$$\nu_n = \nu_0 + n$$

$$\nu_n \sigma_n^2 = \nu_0 \sigma_0^2 + (n - 1)s^2 + \frac{\kappa_0 n}{\kappa_0 + n} (\bar{x} - \mu_0)^2$$

$$\text{where } \bar{x} = \sum_{i=1}^n x_i \text{ and } (n - 1)s^2 = \sum_{i=1}^n (x_i - \bar{x})^2$$

Figure 49: Prior-to-Posterior updating for the iid Gaussian model with unknown mean and variance using the conjugate prior.

- the reason why σ^2 does not appear in w is that the prior variance for θ is scaled by σ^2 in the conjugate prior, and σ^2 therefore cancels out in w .
- the posterior sample size κ_n is the number of prior observations κ_0 plus the sample size n .

Interest centers mainly on the average download speed, so we would like to obtain the marginal posterior distribution of θ . This distribution can be derived by marginalizing out the nuisance parameter σ^2 from the joint posterior

$$p(\theta | x_1, \dots, x_n) = \int p(\theta | \sigma^2, x_1, \dots, x_n) p(\sigma^2 | x_1, \dots, x_n) d\sigma^2,$$

where $p(\theta | \sigma^2, x_1, \dots, x_n)$ and $p(\sigma^2 | x_1, \dots, x_n)$ are given in Figure 49. In Exercise 1 you are asked to show that the marginal posterior of θ is a student- t distribution; see Figure 41 and 42 for a definition and properties. Specifically, we have the following result

$$\theta | x_1, \dots, x_n \sim t(\mu_n, \sigma_n^2 / \kappa_n, \nu_n), \quad (30)$$

where μ_n , σ_n^2 , κ_n and ν_n are all defined as in Figure 49. Note that also the marginal prior for θ follows a student- t distribution of the form (30), but with hyperparameters naturally subscripted by 0 instead of n .

EXAMPLE: INTERNET SPEED DATA. Let us return to the example with the $n = 5$ download speeds with a mean of $\bar{x} = 15.998$ Mbit/s from the chapter [Single-parameter models](#). This time we assume also σ^2 , the variability of the measurements from the speed testing service, to

be unknown. I will use the prior hyperparameters $\mu_0 = 20$, $\kappa_0 = 1$, $\nu_0 = 5$ and $\sigma_0^2 = 5^2$, which agrees in location with my previous prior when σ^2 was assumed known at $\sigma^2 = 5^2$; setting $\nu_0 = 5$ gives a prior equal to the green distribution in the right graph of Figure 51, which I find sensible.




Figure 50: Prior-to-Posterior updating for the internet speed data in the iid Normal model. Contours of joint distributions of θ and σ^2 .

Figure 50 displays contours of the joint prior, likelihood and posterior for θ and σ^2 ; the posterior is more concentrated than the prior, especially for θ . The marginal priors and posterior for the two parameters are shown in Figure 51. The data have made both marginal posteriors more concentrated, but less so for σ^2 since we do not learn so much about a variance from only $n = 5$ observations. The probability of at least 20 Mbit download speed has decreased from the prior probability of 0.5 to 0.066 in the posterior.




Figure 51: Marginal posteriors for the internet speed data in the iid Normal model.

A first look at Monte Carlo simulation

The iid Gaussian model with conjugate prior is an example of a model where we can obtain both the joint and the marginal pos-

teriors in analytical form. This will seldom be the case in more complex models or when non-conjugate priors are used. The idea with Monte Carlo methods is to simulate **posterior draws** of θ from $p(\theta|x_1, \dots, x_n)$ and approximate the posterior by for example a histogram. We will have much more to say about this in Chapter [Posterior simulation](#) where powerful simulation algorithms are presented, but we will already here introduce the most basic Monte Carlo simulation method.

The algorithm in Figure 52 gives pseudo-code for simulating from the $p(\theta, \sigma^2|\mathbf{x})$ in the iid normal model by iteratively simulating from $p(\sigma^2|\mathbf{x})$ followed by simulation from $p(\theta|\sigma^2, \mathbf{x})$. Note how this involves using the most recently simulated value of σ^2 when simulating θ . The algorithm includes the subfunction $\text{RINVCHI2}(v_n, \sigma_n^2)$ to draw from the $\text{Inv-}\chi^2$ distribution. The algorithm implicitly assumes that the standard library of your programming language includes random number generators $\text{RCHI2}(v)$ and $\text{RNORMAL}(\mu_n, \sigma^2/\kappa_n)$ for the χ^2 and normal distributions, respectively.

posterior draws

Posterior simulation - iid Gaussian with conjugate prior.

Input: data $\mathbf{x} = (x_1, \dots, x_n)$
 number of posterior draws m .
 compute $\mu_n, \sigma_n^2, \kappa_n$ and v_n using Figure 49.
for i in $1:m$ **do**
 $\sigma^2 \leftarrow \text{RINVCHI2}(v_n, \sigma_n^2)$
 $\theta \leftarrow \text{RNORMAL}(\mu_n, \sigma^2/\kappa_n)$
end
Output: m draws for θ and σ^2 from joint posterior.

Function $\text{RINVCHI2}(v, \tau^2)$
 $x = \text{RCHI2}(v)$
 $y = v\tau^2/x$
 return y

Figure 52: Algorithm for posterior simulation for the iid Normal model with conjugate prior. The RNORMAL and RCHI2 random number generators are assumed to be part of the standard library. The variable σ^2 is highlighted in orange to indicate that the most recent draw of σ^2 is used in the call to the RNORMAL function.

EXAMPLE: INTERNET SPEED DATA


Let us now use the algorithm in Figure 52 to simulate from the posterior of θ and σ^2 in the Internet speed data. The second and third columns in Table 1 shows the output from generating $m = 10,000$ joint posterior draws with the algorithm in Figure 52. One attractive feature of simulating from the joint posterior distribution is that all marginal posterior distributions are directly obtained by just selecting the column for the parameter in question; tedious in-

tegration is replaced by plotting a histogram of the selected column. Figure 53 shows the marginals obtained from simulation with the analytical marginal posteriors, which happen to be known in this simple example.

The histograms of the simulated draws in Figure 53 are clearly approximating the posteriors extremely well. Monte Carlo simulation is theoretically known to be **simulation consistent** in the sense that we are guaranteed to get arbitrary close to the true posterior if we simulate a large number of draws. For example, if we let $\theta^{(i)}$ denote the i th posterior draw of any of the parameters in a model, then the law of large numbers implies

$$\bar{\theta}_{1:m} \equiv \frac{1}{m} \sum_{i=1}^m \theta^{(i)} \xrightarrow{a.s.} \mathbb{E}(\theta|\mathbf{x}) \text{ as } m \rightarrow \infty,$$

where $\xrightarrow{a.s.}$ denotes almost sure convergence, the strongest form of probabilistic convergence. The result says that the mean of the posterior draws will get closer and closer to the theoretical posterior mean $\mathbb{E}(\theta|\mathbf{x})$ as we increase the number of simulations, m . Figure 54 illustrates this convergence by plotting the posterior mean estimates $\bar{\theta}_{1:m}$ for increasing m ; note that the figure shows the cumulative estimates only up to $m = 1000$.



draw	θ	σ^2
1	18.165	18.451
2	20.431	29.943
3	15.565	29.094
\vdots	\vdots	\vdots
10,000	16.400	21.668
Mean	16.645	30.813

Table 1: Posterior simulation output for the Internet speed dataset.

simulation consistent

Figure 53: Histogram of simulated marginal posteriors for the internet speed data with analytical marginal posterior densities overlaid.

The central limit theorem (CLT) can be used to prove that $\bar{\theta}_{1:m}$ converges in distribution to a normal distribution. Hence, the following approximation of the posterior estimate $\bar{\theta}_{1:m}$ is accurate when m is large:

$$\bar{\theta}_{1:m} \sim N\left(\mathbb{E}(\theta|\mathbf{x}), \frac{\mathbb{V}(\theta|\mathbf{x})}{m}\right), \quad (31)$$

where $\mathbb{V}(\theta|\mathbf{x})$ is the posterior variance of θ ; note that we get the usual reduction in variance that comes from taking averages of m draws, i.e. the variance of $\bar{\theta}_{1:m}$ decreases with m . The result in (31) can be used to determine the required number of draws m needed for a given estimation precision. A multivariate version of the CLT can be used to prove a similar result to (31) when θ is a vector; an interest-




Figure 54: Convergence of the Monte Carlo estimate of the posterior expectation of θ and σ^2 . The analytical posterior expectations are displayed as thin horizontal lines.

ing aspect is that $\text{Cov}(\bar{\theta}_{1:m})$ (a covariance matrix in the multiparameter case) still decreases at the rate $1/m$, regardless of the dimension of θ .




Figure 55: Histogram of simulated marginal posteriors for σ (left) and the coefficient of variation σ/θ (right) for the internet speed data.

It is often the case that the quantities of interest are functions $f(\theta)$ of the parameters; for example the **coefficient of variation** σ/θ in the iid normal model. Even when the posterior for the model parameters θ is available analytically, deriving the posterior for $f(\theta)$ involves tedious multidimensional change-of-variables calculations. Here is a second attractive property of simulation: the posterior for $f(\theta)$ can be directly obtained from a posterior sample of θ by simply computing the function $f(\theta)$ for each posterior draw. Provided the posterior variance of $f(\theta)$ exists, a central limit theorem of the form (31) exists also in this case, with the expected value and variance replaced by those of $f(\theta)$.

To illustrate how simulation immediately provides inference for any function of the parameters, Table 2 contains a fourth column named σ/θ with the computed coefficient of variation for each draw. We can now just plot a histogram of this new column to approximate the marginal posterior of the function $f(\theta, \sigma^2) = \sigma/\theta$. The results are presented in the right part of Figure 55; the left part of the figure shows the results for the standard deviation $f(\theta, \sigma^2) = \sqrt{\sigma^2}$.

The final column of Table 2 is a binary variable that records if θ was at least 20, i.e. it computes the indicator function $f(\theta, \sigma^2) = I(\theta \geq 20)$. The marginal posterior probability $\Pr(\theta \geq 20|\mathbf{x})$ is then easily approximated by the mean of the final column; Figure 56 illustrates the Monte Carlo convergence of this estimate.

Multinomial data

Categorical data have observations that belong to one of C discrete classes. A computer bug can for example be allocated to C developing teams; an items sold in an auction may reported as: 'defective', 'normal quality', or 'new'; a continuous variable like age can recorded in age intervals: 0-18, 19-28, 29-49, 50-64 and 65+, which

draw	θ	σ^2	σ/θ	$\theta \geq 20$
1	18.165	18.451	0.236	0
2	20.431	29.943	0.267	1
3	15.565	29.094	0.346	0
\vdots	\vdots	\vdots	\vdots	\vdots
10,000	16.400	21.668	0.283	0
Mean	16.645	30.813	0.330	0.066

Table 2: Posterior simulation output for the Internet speed dataset with computed functions of the parameters.

coefficient of variation




Figure 56: Convergence of the Monte Carlo estimate of $\Pr(\theta \geq 20|\mathbf{x})$. The analytical posterior probability is displayed as a thin horizontal line.

Categorical data

would then also be a categorical variable. The categories in the latter two situations are examples of **ordinal data** where the categories have a natural order. There are special models for ordinal data which we will not cover in this chapter; here we will consider categorical data without natural order. Categorical variables are often called **multi-class** in the machine learning literature.

A multi-class random variable X is often written in **one-hot encoding** as $\mathbf{x} = (x_1, \dots, x_C)$ where $X = c$ is encoded as $x_c = 1$ and $x_j = 0$ for $j \neq c$; hence when $C = 3$, $\mathbf{x} = (0, 1, 0)$ means that the observation belongs to the second class. The categorical random variable $X|\boldsymbol{\theta} \sim \text{Cat}(\theta_1, \dots, \theta_C)$ has probability distribution

$$p(\mathbf{x}) = \theta_1^{x_1} \dots \theta_C^{x_C}, \quad (32)$$

where (x_1, \dots, x_C) is the one-hot encoding of x , $0 < \theta_c < 1$ is the probability of class c and $\sum_{c=1}^C \theta_c = 1$. Note how Bernoulli data is the special case with $C = 2$ categories 'success' and 'failure', so that the $\text{Cat}(\theta_1, \dots, \theta_C)$ distribution generalizes the Bernoulli distribution to the case $C > 2$. Figure 57 is an example of $\text{Cat}(\theta_1, \dots, \theta_C)$ for $C = 4$.

We saw in Section [The likelihood function and maximum likelihood estimation](#) that counting the number of successes s in n binary Bernoulli trials gave rise to $S \sim \text{Binomial}(n, \theta)$ data. In the same way we can count the number of observations in category c for $c = 1, \dots, C$ in multi-class data. This gives data as a count vector $\mathbf{y} = (y_1, \dots, y_C)$ where y_c is the number of observations in category c in $n = \sum_{c=1}^C y_c$ 'trials'. Here is an example:

SMARTPHONE SURVEY DATA. A company conducted a survey among $n = 513$ smartphone users. Among other questions, the participants were asked: 'What kind of mobile phone do you mainly use?' with the four options: i) iPhone, ii) Android, iii) Windows and iv) Other/Don't know. The number of responses in the four categories were: $\mathbf{y} = (180, 230, 62, 41)$.

The **multinomial distribution** generalizes the binomial distribution to $C > 2$ categories; its main properties are summarized in Figure 58. The Binomial distribution in Figure 6 is the special case with $C = 2$ categories, which is seen by defining $\theta = \theta_1$, $\theta_2 = 1 - \theta$, $x = x_1$, $x_2 = n - x$, and noting that

$$\frac{n!}{x_1!x_2!} = \frac{n!}{x!(n-x)!} = \binom{n}{x}. \quad (33)$$

The multinomial distribution is a multivariate distribution with convenient marginalization properties. For example, if we group the counts in one or more categories - for example turning the smartphone dataset into three categories by merging 'Windows' and

ordinal data

multi-class

one-hot encoding




Figure 57: Categorical distribution with probabilities $\boldsymbol{\theta} = (0.20, 0.30, 0.05, 0.45)$.

Multinomial distribution

$(X_1, \dots, X_C) \sim \text{MultiNom}(n, \boldsymbol{\theta})$
 where $\sum_{c=1}^C X_c = n$,
 $\boldsymbol{\theta} = (\theta_1, \dots, \theta_C)$ and $\sum_c \theta_c = 1$.

$$p(\mathbf{x}) = \frac{n!}{x_1! \dots x_C!} \theta_1^{x_1} \dots \theta_C^{x_C}$$

$$\mathbb{E}(X_c) = n\theta_c$$

$$\mathbb{V}(X_c) = n\theta_c(1 - \theta_c)$$

Figure 58: The multinomial distribution.

multinomial distribution

'Other' - the distribution remains multinomial. The probability of a merged category is simply the sum of the probabilities of the merged categories. Hence

$$(x_1, x_2, x_3 + x_4) \sim \text{Multinomial}(\theta_1, \theta_2, \theta_3 + \theta_4).$$

In particular, merging to only two categories - for example 'iPhone' and 'not iPhone' - gives a binomial distribution where the probability of failure (not iPhone) is $\theta_2 + \theta_3 + \theta_4$.

A Bayesian analysis of multinomial data requires a prior distribution for the model parameters, $\theta = (\theta_1, \dots, \theta_C)$. Since each θ_c is a probability, the first distribution that comes to mind may be a Beta distribution; the Beta distribution is not appropriate here however since it does not enforce the constraint that the probabilities sum to one. Hence, the parameter space of the multinomial distribution is the **unit simplex**, i.e. the set $\theta = (\theta_1, \dots, \theta_C) : 0 < \theta_c < 1$ and $\sum_c \theta_c = 1$. Luckily, there is a very nice distribution on the unit simplex, the Dirichlet distribution, summarized in Figure 59.

The Dirichlet distribution is specified with the prior hyperparameters $\alpha_c > 0$, see Figure 60 for some examples. The *relative* sizes of the elements in α determine the prior means for elements of θ . For example, setting $\alpha_1 = \dots = \alpha_C = 1.5$, as in the upper graph of Figure 60, gives equal prior mean for all categories: $\mathbb{E}(\theta_c) = 1/C$ for all c . The *absolute* size of α , measured by $\alpha_+ = \sum_{c=1}^C \alpha_c$, is inversely related to the variance, see Figure 59; hence, the prior hyperparameters $\alpha = (1.5, \dots, 1.5)$ and $\alpha = (5, \dots, 5)$ in the upper and middle part of Figure 60 have the same mean, but the latter has smaller variance. Finally, the bottom part of Figure 60 shows an example where the prior mean is different over the categories.

The Dirichlet(1, ..., 1) has constant density and is therefore the **uniform distribution on the unit simplex**; this generalizes the result that Beta(1, 1) is uniform on the unit interval [0, 1]. Finally, when $\alpha_c < 1$, the Dirichlet density becomes 'bathtub shaped' with probability mass piling up against the edges of the unit simplex.

The Dirichlet distribution is conjugate to the multinomial likelihood which is easily seen by computing the posterior

$$p(\theta|\mathbf{y}) \propto p(\mathbf{y}|\theta)p(\theta) \quad (34)$$

$$= \frac{n!}{x_1! \dots x_C!} \theta_1^{x_1} \dots \theta_C^{x_C} \cdot \frac{\Gamma(\sum_{c=1}^C \alpha_c)}{\prod_{c=1}^C \Gamma(\alpha_c - 1)} \theta_1^{\alpha_1-1} \dots \theta_C^{\alpha_C-1} \quad (35)$$

$$= \theta_1^{\alpha_1+x_1-1} \dots \theta_C^{\alpha_C+x_C-1}, \quad (36)$$

which is proportional to the Dirichlet($\alpha_1 + x_1, \dots, \alpha_C + x_C$) density. This is a convenient result: the posterior is simply obtained by adding the data count x_c to the prior hyperparameter α_c in each

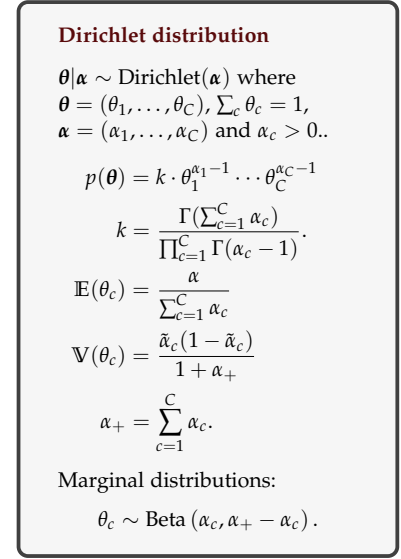


Figure 59: The Dirichlet distribution.




Figure 60: Examples of Dirichlet distributions for $\mathbf{x} = (x_1, x_2, x_3)$.

Source: Wikipedia.

unit simplex

uniform distribution on the unit simplex

category. This parallels and generalizes the binary case where a $\text{Beta}(\alpha, \beta)$ prior was updated to a posterior by adding the number of successes s to α and the number of failures f to β . Figure 61 summarizes the prior-to-posterior updating for multinomial data with a Dirichlet prior.

Multinomial data with Dirichlet prior

Model: $\mathbf{n}|\boldsymbol{\theta} \sim \text{Multinomial}(\boldsymbol{\theta})$, where
 $\mathbf{n} = (n_1, \dots, n_C)$ are counts in C categories
 $\boldsymbol{\theta} = (\theta_1, \dots, \theta_C)$ are category probabilities.
Prior: $\boldsymbol{\theta} \sim \text{Dirichlet}(\boldsymbol{\alpha})$, for $\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_C)$
Posterior: $\boldsymbol{\theta} \sim \text{Dirichlet}(\boldsymbol{\alpha} + \mathbf{n})$

Figure 61: Prior-to-Posterior updating for multinomial data with the Dirichlet prior.

SMARTPHONE SURVEY DATA We are now ready to analyze the four market shares $\theta_1, \dots, \theta_4$ in the smartphone data. We will determine the prior hyperparameters in the Dirichlet prior using data from a similar survey from four year ago. The proportions in the four categories back then were: 30%, 30%, 20% and 20%. This was a large survey, but since time has passed and user patterns most likely have changed, I value the information in this older survey as being equivalent to a survey with only 50 participants. This gives us the prior:

$$(\theta_1, \dots, \theta_4) \sim \text{Dirichlet}(\alpha_1 = 15, \alpha_2 = 15, \alpha_3 = 10, \alpha_4 = 10)$$

Note that $\mathbb{E}(\theta_1) = 15/50 = 0.3$ and so on, so the prior mean is set equal to the proportions from the older survey. Also, $\sum_{k=1}^4 \alpha_k = 50$, so the prior information is equivalent to a survey based on 50 respondents, as required.

The joint posterior distribution of all four shares is by Figure 61 equal to

$$(\theta_1, \dots, \theta_4) | \mathbf{y} \sim \text{Dirichlet}(15 + 180, 15 + 230, 10 + 62, 10 + 41)$$

The marginal posteriors are plotted in Figure 62 as histograms from Monte Carlo simulation (see the algorithm in Figure 63); the analytical posteriors from Figure 59 are overlaid.

Figure 62 indicates that Android may have the largest market share with a posterior mean around 0.44 versus iPhones posterior mean of 0.35. Computing the probability that Android has the largest market share involves integrating the joint posterior $\boldsymbol{\theta} | \mathbf{y} \sim \text{Dirichlet}(\boldsymbol{\alpha} + \mathbf{y})$ over the region $\{\boldsymbol{\theta} : \theta_2 > \max(\theta_1, \theta_3, \theta_4)\}$, a tedious calculation. The probability is however easily computed by

draw	θ_1	θ_2	θ_3	θ_4	I
1	0.33	0.47	0.10	0.09	1
2	0.34	0.44	0.11	0.09	1
3	0.36	0.41	0.13	0.08	1
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
10,000	0.35	0.43	0.14	0.08	1
Mean	0.34	0.43	0.13	0.09	0.99

Table 3: Posterior simulation output for the multinomial model applied to the Smartphone survey data. The last column is a computed binary indicator for the event that Android has the largest market share, i.e. if $\theta_2 > \max(\theta_1, \theta_3, \theta_4)$.

simulation by recording for each posterior θ draw if the condition $\theta_2 > \max(\theta_1, \theta_3, \theta_4)$ is satisfied; Table 3 shows that

$$\Pr(\text{Android has largest market share} | \mathbf{y}) \approx 0.99,$$

so we can be almost certain that Android is the most popular smartphone operating system.




Figure 62: Marginal posteriors of the market shares for the smartphone survey data. Simulated (histogram) draws and analytical density functions (solid curves).

Posterior simulation - Multinomial data, Dirichlet prior.

Input: data $\mathbf{n} = (n_1, \dots, n_C)$
 prior hyperparameters $\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_C)$
 the number of posterior draws m .

for i in $1:m$ **do**
 | $\theta \leftarrow \text{RDIRICHLET}(\boldsymbol{\alpha} + \mathbf{n})$

end

Output: m posterior draws of $\theta = (\theta_1, \dots, \theta_C)$.

Function $\text{RDIRICHLET}(\boldsymbol{\alpha})$

for c in $1:C$ **do**
 | $y[c] \leftarrow \text{RGAMMA}(\alpha[c], 1)$
end
return $\mathbf{y} / \text{SUM}(\mathbf{y})$

Figure 63: Algorithm for posterior simulation for the multinomial model with the conjugate Dirichlet prior. The `RGAMMA` random number generator is assumed to be part of the standard library.

Prior elicitation and objective Bayes

Noninformative priors

- Flat - No. prior data -> 0

Jeffreys' prior

- Jeffreys prior. Do also $1/\sigma^2$ thing here. - Reference priors - the idea

*Prior elicitation***EXERCISES**

1. Derive the marginal posterior of θ in (30) for the iid Gaussian model $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} N(\theta, \sigma^2)$.
2. Let $x_1, \dots, x_n | \theta \stackrel{\text{iid}}{\sim} N(\theta, \sigma^2)$, where θ is assumed known. Show that the $\text{Inv-}\chi^2$ distribution is a conjugate prior for σ^2 .

NOTEBOOK EXERCISES

1. See the notebook [multi-parameter models](#).

Regression

Regression models are the most important of all statistical models as they appear as a component in nearly any situation where an output variable y is modeled as function of a set of input variables $\mathbf{x} = (x_1, \dots, x_p)^\top$, where \top denotes vector transpose. The input variables are often called **covariates**, predictors or **features**, and the output variable is most commonly termed the **response variable** or target variable. In the chapter [Classification](#) we will see regression models for a binary response variable and also for response variables of other data types, for example counts. Regression is also the basis for deep neural networks where a linear combination of covariates are passed through several nonlinear activation functions before finally being linked to the response.

The basic **Gaussian linear regression model** is

$$y_i = \mathbf{x}_i^\top \boldsymbol{\beta} + \varepsilon_i, \quad \varepsilon_i \stackrel{iid}{\sim} N(0, \sigma^2) \quad \text{for } i = 1, \dots, n,$$

where \mathbf{x}_i is a vector with observations on the p covariates and $\boldsymbol{\beta} = (\beta_1, \dots, \beta_p)^\top$ is the vector of **regression coefficients**. The β_j are called **weights** in the machine learning literature and are therefore frequently denoted by w_j . The model is said to be **homoscedastic** since the error variance σ^2 is the same for all observations.

It is convenient to stack all n response observations in a vector $\mathbf{y} = (y_1, \dots, y_n)^\top$ and the covariate observations vectors as rows in the $n \times p$ covariate matrix $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)^\top$. The Gaussian linear regression model can then be expressed as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \stackrel{iid}{\sim} N(0, \sigma^2 I_n), \quad (37)$$

where $\boldsymbol{\varepsilon}$ is vector with all the ε_i and $N(0, \sigma^2 I_n)$ is the multivariate normal distribution with diagonal covariance matrix $\sigma^2 I_n$ and I_n is the identity matrix; the simple diagonal structure of $\text{Cov}(\boldsymbol{\varepsilon})$ reflects the assumption that the ε_i are independent with the same variance.

Likelihood and MLE

The likelihood for the linear regression model with homoscedastic Gaussian errors is given by the following multivariate normal distri-

covariates
features
response variable

Gaussian linear regression
model

regression coefficients
weights
homoscedastic

bution

$$\mathbf{y}|\beta, \sigma^2, \mathbf{X} \sim N(\mathbf{X}\beta, \sigma^2 I_n), \quad (38)$$

where we note that the covariates \mathbf{X} are assumed fixed so the likelihood is the distribution of only the response \mathbf{y} .

The **least squares estimator** $\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ is well known to minimize the sum of squared **residuals**

$$Q(\beta) \equiv (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta).$$

When the errors are homoscedastic Gaussian, $\hat{\beta}$ is also the MLE since the log-likelihood from (38) is a constant plus $-(1/2\sigma^2)Q(\beta)$.

The sampling distribution of the MLE is easily obtained since $\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ is a linear function of \mathbf{y} and $(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$ is a constant matrix. Since $\mathbf{y}|\beta, \sigma^2, \mathbf{X} \sim N(\mathbf{X}\beta, \sigma^2 I_n)$, the frequentist sampling distribution of $\hat{\beta}$ is obtained by applying the result in Figure 64 with $\mu = \mathbf{X}\beta$, $\Sigma = \sigma^2 I_n$ and $\mathbf{A} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$ to obtain

$$\hat{\beta}|\beta, \sigma^2, \mathbf{X} \sim N(\beta, \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1}),$$

from which we also see that the MLE is unbiased for β .

The MLE for σ^2 can be shown to be $\hat{\sigma}^2 \equiv (\mathbf{y} - \mathbf{X}\hat{\beta})^\top (\mathbf{y} - \mathbf{X}\hat{\beta})/n$. The estimator $\hat{\sigma}^2$ is biased for σ^2 , and the following unbiased estimator is typically used instead

$$s^2 \equiv \frac{(\mathbf{y} - \mathbf{X}\hat{\beta})^\top (\mathbf{y} - \mathbf{X}\hat{\beta})}{n - p}.$$

Non-informative prior

We will start with a noninformative improper prior which is uniform over β and $\phi \equiv \log \sigma^2$. With a uniform prior on $\log \sigma^2$ we can calculate the implied prior on σ^2 by the change-of-variable formula as

$$p(\sigma^2) \propto p_\phi(\log \sigma^2) \left| \frac{d \log \sigma^2}{d \sigma^2} \right| \propto \frac{1}{\sigma^2},$$

since $p_\phi(\cdot)$ is uniform, i.e. proportional to a constant.

The joint posterior for β and σ^2 is given by Bayes' theorem as

$$\begin{aligned} p(\beta, \sigma^2 | \mathbf{y}) &\propto p(\mathbf{y} | \beta, \sigma^2) p(\beta, \sigma^2) \propto N(\mathbf{y} | \mathbf{X}\beta, \sigma^2 I_n) \cdot \frac{1}{\sigma^2} \\ &= |2\pi\sigma^2 I_n|^{-1/2} \exp \left\{ -\frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta) \right\} \cdot \frac{1}{\sigma^2}, \end{aligned} \quad (39)$$

where the conditioning on the fixed covariates \mathbf{X} is suppressed to simplify the notation. Now, $(\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)$ can be rewritten using the MLE $\hat{\beta}$ as

$$(\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta) = (\mathbf{y} - \mathbf{X}\hat{\beta})^\top (\mathbf{y} - \mathbf{X}\hat{\beta}) + (\beta - \hat{\beta})^\top \mathbf{X}^\top \mathbf{X} (\beta - \hat{\beta}), \quad (40)$$

least squares estimator
residuals

Linear transformation of Gaussians

Let $\mathbf{x} \sim N(\mu, \Sigma)$ be multivariate Gaussian and \mathbf{A} a constant full rank $m \times p$ matrix. Then

$$\mathbf{A}\mathbf{x} \sim N(\mathbf{A}\mu, \mathbf{A}\Sigma\mathbf{A}^\top).$$

Figure 64: Linear transformation of Gaussians.

which can be directly verified by substituting the definition of $\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$. Recall from linear algebra that the determinant of a diagonal matrix is the product of its diagonal elements, so $|2\pi\sigma^2 I_n| = (2\pi\sigma^2)^n \propto (\sigma^2)^n$. Using this result and (40) in (39) we obtain the posterior

$$p(\beta, \sigma^2 | \mathbf{y}) \propto (\sigma^2)^{-(n+2)/2} \exp \left\{ -\frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\hat{\beta})^\top (\mathbf{y} - \mathbf{X}\hat{\beta}) \right\} \quad (41)$$

$$\cdot \exp \left\{ -\frac{1}{2\sigma^2} (\beta - \hat{\beta})^\top \mathbf{X}^\top \mathbf{X} (\beta - \hat{\beta}) \right\} \quad (42)$$

The posterior is most transparent if we use the decomposition of the joint posterior

$$p(\beta, \sigma^2 | \mathbf{y}) = p(\beta | \sigma^2, \mathbf{y}) p(\sigma^2 | \mathbf{y}).$$

Focusing first on $p(\beta | \sigma^2, \mathbf{y}, \mathbf{X})$ we only need to be concerned with the last factor in (41) as it is the only part that depends on β ; note that $\hat{\beta}$ only depends on the data. We immediately recognize this last factor as proportional to the multivariate normal density, so

$$\beta | \sigma^2, \mathbf{y} \sim N(\hat{\beta}, \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1}).$$

The marginal posterior of σ^2 is obtained by integrating out β in (41)

$$\begin{aligned} p(\sigma^2 | \mathbf{y}) &= \int p(\beta, \sigma^2 | \mathbf{y}) d\beta \\ &\propto (\sigma^2)^{-(n+2)/2} \exp \left\{ -\frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\hat{\beta})^\top (\mathbf{y} - \mathbf{X}\hat{\beta}) \right\} \\ &\quad \cdot \int \exp \left\{ -\frac{1}{2\sigma^2} (\beta - \hat{\beta})^\top \mathbf{X}^\top \mathbf{X} (\beta - \hat{\beta}) \right\} d\beta \\ &\propto (\sigma^2)^{-(n+2)/2} \exp \left\{ -\frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\hat{\beta})^\top (\mathbf{y} - \mathbf{X}\hat{\beta}) \right\} (\sigma^2)^{p/2}, \end{aligned}$$

where the last proportionality comes from the fact that

$$\int \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\} d\mathbf{x} = |2\pi\boldsymbol{\Sigma}|^{1/2}$$

for any p -vectors \mathbf{x} and $\boldsymbol{\mu}$, and positive definite matrix $\boldsymbol{\Sigma}$ since we know that the $N(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Sigma})$ density integrates to one over \mathbb{R}^p . The marginal posterior for σ^2 is therefore

$$p(\sigma^2 | \mathbf{y}) \propto (\sigma^2)^{-[1+(n-p)/2]} \exp \left\{ -\frac{1}{2\sigma^2} (n-p)s^2 \right\}, \quad (43)$$

which can be recognized as proportional to the $\text{Inv-}\chi^2(n-p, s^2)$ density.

We summarize the prior-to-posterior updating in Gaussian linear regression with a noninformative prior in Figure 65.

Gaussian linear regression with non-informative prior**Model:** $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \stackrel{iid}{\sim} \mathcal{N}(0, \sigma^2 \mathbf{I}_n)$ **Prior:** $p(\boldsymbol{\beta}, \sigma^2) \propto 1/\sigma^2$ **Posterior:** $\boldsymbol{\beta} | \sigma^2, \mathbf{y}, \mathbf{X} \sim \mathcal{N}(\hat{\boldsymbol{\beta}}, \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1})$
 $\sigma^2 | \mathbf{y}, \mathbf{X} \sim \text{Inv-}\chi^2(n-p, s^2)$ where $\hat{\boldsymbol{\beta}} \equiv (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ and $s^2 \equiv (\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}})^\top (\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}) / (n-p)$.

Figure 65: Prior-to-Posterior updating for the Gaussian linear regression with non-informative prior.

Conjugate prior

Let us now turn to the more interesting case with a conjugate prior for the Gaussian linear regression. Recall the form of the conjugate prior for the iid normal model

Gaussian linear regression with conjugate prior**Model:** $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \stackrel{iid}{\sim} \mathcal{N}(0, \sigma^2 \mathbf{I}_n)$ **Prior:** $\boldsymbol{\beta} | \sigma^2 \sim \mathcal{N}(\boldsymbol{\mu}_0, \sigma^2 \boldsymbol{\Omega}_0^{-1})$
 $\sigma^2 | \mathbf{y}, \mathbf{X} \sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2)$ **Posterior:** $\boldsymbol{\beta} | \sigma^2, \mathbf{y}, \mathbf{X} \sim \mathcal{N}(\hat{\boldsymbol{\beta}}, \sigma^2 (\mathbf{X}^\top \mathbf{X})^{-1})$
 $\sigma^2 | \mathbf{y}, \mathbf{X} \sim \text{Inv-}\chi^2(\nu_n, \sigma_n^2)$ where $\hat{\boldsymbol{\beta}} \equiv (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ and $s^2 \equiv (\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}})^\top (\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}) / (n-p)$.

Figure 66: Prior-to-Posterior updating for the Gaussian linear regression with non-informative prior.

EXERCISES

1. This is the first problem.
2. This is the second problem.

NOTEBOOK EXERCISES

1. See the notebook [regression](#).

Prediction and Decision making

EXERCISES

1. (a) Let $x_1, \dots, x_n \stackrel{iid}{\sim} \text{Bern}(\theta)$, with a $\text{Beta}(\alpha, \beta)$ prior for θ . Derive the predictive distribution for x_{n+1} .
(b) You need to decide if you bring your umbrella during your daily walk. It has rained on two days during the last ten days, and you assess those ten days to be representative also for the weather today, the 11th day. Your utility for the action-state combinations are given in the table below. Assume a $\text{Beta}(1, 1)$ prior for θ . Compute the Bayesian decision.
(c) How sensitive is your decision in (b) to the changes in the prior hyperparameters, α and β ?
2. (a) Let x_i be the number of sales of a product on month i . Let $x_1, \dots, x_n \stackrel{iid}{\sim} N(\theta, \sigma^2)$ be the (approximate) distribution for the sales, and let $\theta \sim N(200, 50^2)$ a priori. Assume that $\sigma^2 = 25^2$ and that we have observed $n = 5$ and $\bar{x} = 320.4$. Compute the predictive distribution for x_6 .
(b) The company has the choice of performing a marketing campaign for their product. The marketing campaign costs 300 and is believed to increase sales by 20% compared to when no campaign is performed. The company sells the product for $p = 10$ dollar and the cost of producing the product is $q = 5$ dollar. There are no fixed production costs. Assume that the company's utility is described by $U(y) = 1 - \exp(-y/1000)$, where y is the total profit from sales in the next month. Should the company perform the marketing campaign? *Hint: the expected value of the exponential function of a normal random variable $S \sim N(\mu, \sigma^2)$ is $\mathbb{E}(\exp(S)) = \exp(\mu + \sigma^2/2)$.*

NOTEBOOK EXERCISES

1. See the notebook [Prediction and Decision](#).

Prediction in normal model with known variance

My streaming service becomes unreliable and buffers at speeds below 5Mbit/sec, and what I really like to know is the probability of this 'catastrophic' event. Finding the probability that a *single* measurement will be lower than 5MBit/sec is instead an exercise in prediction.

whatever

- Exchangeability and De Finettis theorem

whatever




Figure 67: WiFi example. Predictive density.

Classification

EXERCISES

1. This is the first problem.
2. This is the second problem.

NOTEBOOK EXERCISES

1. See the notebook [Classification](#).

Posterior simulation

Gibbs sampling

Markov Chain Monte Carlo

Hamiltonian Monte Carlo

Probabilistic programming frameworks

Variational inference

Regularization

Model comparison

Posterior model probabilities

Bayesian cross-validation

- M-completed? - Generalization performance. - WAIC - Cross-validation This is some text ³

L2-regularization and Ridge

L1-regularization and Lasso

Global-local regularization and Horseshoe

³ Andrew Gelman, John B Carlin, Hal S Stern, David B Dunson, Aki Vehtari, and Donald B Rubin. *Bayesian data analysis*, volume 3rd edition. CRC press, 2013

Variable selection

Gaussian processes

Gaussian processes

Mixture models

Finite mixtures

Mixtures of regressions

Latent Dirichlet allocation

Infinite mixtures

Bibliography

José M Bernardo and Adrian FM Smith. *Bayesian theory*, volume 405. John Wiley & Sons, 2009.

George Casella and Roger L Berger. *Statistical inference*, volume 2. Duxbury Pacific Grove, CA, 2002.

Dheeru Dua and Casey Graff. UCI machine learning repository, 2017.
URL <http://archive.ics.uci.edu/ml/datasets/Spambase/>.

Andrew Gelman, John B Carlin, Hal S Stern, David B Dunson, Aki Vehtari, and Donald B Rubin. *Bayesian data analysis*, volume 3rd edition. CRC press, 2013.

Rolf Sundberg. *Statistical modelling by exponential families*, volume 12. Cambridge University Press, 2019.

Bertil Wegmann and Mattias Villani. Bayesian inference in structural second-price common value auctions. *Journal of Business & Economic Statistics*, 29(3):382–396, 2011.

Index

- batch learning, 26
- Bayes estimator, 19
- Bayes' theorem, 12
- Bernoulli distribution, 8
- Bernoulli trials, 8
- Beta distribution, 17
- Binomial distribution, 10
- Birnbaum's theorem, 22

- Categorical data, 44
- coefficient of variation, 44
- conjugate prior, 18
- covariates, 51
- credibility interval, 30

- dutch book argument, 12

- eBayCoin dataset, 27
- equal tail credibility interval, 30
- exponential family, 32

- Factorization criterion, 31
- features, 51
- frequentist probability, 11

- Gamma distribution, 27
- Gaussian linear regression model, 51

- Highest Posterior Density (HPD)
 - region, 30
- homoscedastic, 51

- hyperparameters, 24

- iid, 7
- imaginary prior sample, 18
- Internet speed dataset, 25, 40, 42
- inverse Gamma distribution, 39

- joint posterior distribution, 37

- law of total probability, 13
- least squares estimator, 52
- license, 2
- likelihood function, 8
- Likelihood principle, 22
- likelihood surface, 37

- maximum likelihood estimator, 9
- multi-class, 45
- multinomial distribution, 45

- natural parameter, 32
- negative binomial distribution, 21
- nuisance parameters, 38

- one-hot encoding, 45
- online learning, 25
- ordinal data, 45

- parameter space, 7
- personal degree of belief, 11
- Poisson distribution, 26

- posterior, 13
- posterior density, 14
- posterior draws, 42
- prior, 13
- prior density, 14
- prior elicitation, 17

- regression coefficients, 51
- residuals, 52
- response variable, 51

- sampling distribution, 10
- sampling variance, 10
- scaled inverse chi-squared distribution, 39
- simulation consistent, 43
- Smartphone survey data, 45
- SpamBase dataset, 19
- student-*t* distribution, 33
- subjective consensus, 15
- subjective probability, 11
- Sufficiency principle, 32
- Sufficient statistic, 31

- unbiased, 9
- uniform distribution, 17
- uniform distribution on the unit simplex, 46
- unit simplex, 46

- weights, 51